

Socket Programming


socket() -- Get the file descriptor

- `int socket(int domain, int type, int protocol);`
 - domain should be set to `AF_INET`
 - type can be `SOCK_STREAM` or `SOCK_DGRAM`
 - set protocol to 0 to have socket choose the correct protocol based on type
 - `socket()` returns a socket descriptor for use in later system calls or -1 on error

```
int sockfd;
```

```
sockfd = socket (PF_INET, SOCK_STREAM, 0);
```

`AF_INET`



Socket Structures

- struct sockaddr: Holds socket address information for many types of sockets

```
struct sockaddr {  
    unsigned short  sa_family;    //address family AF_XXX  
    unsigned short  sa_data[14]; //14 bytes of protocol addr  
}
```

- struct sockaddr_in: A parallel structure that makes it easy to reference elements of the socket address

```
struct sockaddr_in {  
    short int          sin_family;    // set to AF_INET  
    unsigned short int sin_port;      // Port number  
    struct in_addr     sin_addr;      // Internet address  
    unsigned char      sin_zero[8];  //set to all zeros  
}
```

- sin_port and sin_addr must be in **Network Byte Order**

Byte Ordering

- Two types of “Byte ordering”
 - Network Byte Order: High-order byte of the number is stored in memory at the lowest address
 - Host Byte Order: Low-order byte of the number is stored in memory at the lowest address
 - Network stack (TCP/IP) expects Network Byte Order
- Conversions:
 - htons() - Host to Network Short
 - htonl() - Host to Network Long
 - ntohs() - Network to Host Short
 - ntohl() - Network to Host Long

Dealing with IP Addresses

- ```
struct in_addr {
 unsigned long s_addr; // that's a 32-bit long, or 4 bytes
};
```
- ```
int inet_aton(const char *cp, struct in_addr *inp);
```

```
struct sockaddr_in my_addr;  
my_addr.sin_family = AF_INET;  
my_addr.sin_port = htons(MYPORT);  
inet_aton("10.0.0.5",&(my_addr.sin_addr));  
memset(&(my_addr.sin_zero),'\0',8);
```

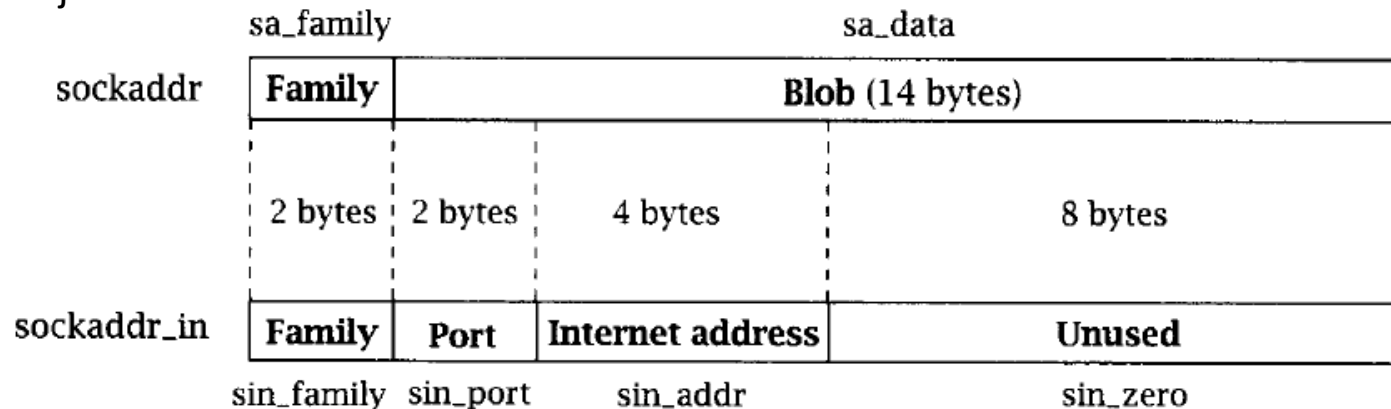
 - `inet_aton()` gives non-zero on success; zero on failure
- To convert binary IP to string: `inet_ntoa()`

```
printf("%s",inet_ntoa(my_addr.sin_addr));
```

❖ Specifying Addresses

- Applications need to be able to specify Internet address and Port number. How?
- Use **Address Structure**
 1. `sockaddr`: generic data type
 2. `in_addr`: internet address
 3. `sockaddr_in`: another view of `sockaddr`

```
struct sockaddr_in{
    unsigned short sin_family; /* Internet protocol (AF_INET) */
    unsigned short sin_port; /* Address port (16 bits) */
    struct in_addr sin_addr; /* Internet address (32 bits) */
    char sin_zero[8]; /* Not used */
}
```



bind() - what port am I on?

- Used to associate a socket with a port on the local machine
 - The port number is used by the kernel to match an incoming packet to a process
- `int bind(int sockfd, struct sockaddr *my_addr, int addrlen)`
 - `sockfd` is the socket descriptor returned by `socket()`
 - `my_addr` is pointer to `struct sockaddr` that contains information about your IP address and port
 - `addrlen` is set to `sizeof(struct sockaddr)`
 - returns -1 on error
 - `my_addr.sin_port = 0; //choose an unused port at random`
 - `my_addr.sin_addr.s_addr = INADDR_ANY; //use my IP adr`

Example

```
int sockfd;  
  
struct sockaddr_in my_addr;  
  
sockfd = socket(PF_INET, SOCK_STREAM, 0);  
  
my_addr.sin_family = AF_INET;      // host byte order  
my_addr.sin_port = htons(MYPORT);  // short, network byte  
                                     order  
  
my_addr.sin_addr.s_addr = inet_addr("172.28.44.57");  
  
memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct  
  
bind(sockfd, (struct sockaddr *)&my_addr, sizeof(struct  
sockaddr));  
  
/***** Code needs error checking. Don't forget to do that *****/
```


connect() - Hello!

- Connects to a remote host
- `int connect(int sockfd, struct sockaddr *serv_addr, int addrlen)`
 - `sockfd` is the socket descriptor returned by `socket()`
 - `serv_addr` is pointer to `struct sockaddr` that contains information on destination IP address and port
 - `addrlen` is set to `sizeof(struct sockaddr)`
 - returns -1 on error
- No need to `bind()`, kernel will choose a port

Example

```
#define DEST_IP  "172.28.44.57"
#define DEST_PORT 5000
main(){
    int sockfd;
    struct sockaddr_in dest_addr; // will hold the destination addr
    sockfd = socket(PF_INET, SOCK_STREAM, 0);
    dest_addr.sin_family = AF_INET; // host byte order
    dest_addr.sin_port = htons(DEST_PORT); // network byte
    order
    dest_addr.sin_addr.s_addr = inet_addr(DEST_IP);
    memset(&(dest_addr.sin_zero), '\0', 8); // zero the rest of the
    struct    connect(sockfd, (struct sockaddr *)&dest_addr,
    sizeof(struct sockaddr));
    /***** Don't forget error checking *****/
}
```

listen() - Call me please!

- Waits for incoming connections
- `int listen(int sockfd, int backlog);`
 - `sockfd` is the socket file descriptor returned by `socket()`
 - `backlog` is the number of connections allowed on the incoming queue
 - `listen()` returns -1 on error
 - Need to call `bind()` before you can `listen()`
 - `socket()`
 - `bind()`
 - `listen()`
 - `accept()`

accept() - Thank you for calling !

- accept() gets the pending connection on the port you are listen()ing on
- `int accept(int sockfd, void *addr, int *addrlen);`
 - sockfd is the listening socket descriptor
 - information about incoming connection is stored in addr which is a pointer to a local struct `sockaddr_in`
 - addrlen is set to `sizeof(struct sockaddr_in)`
 - accept returns *a new socket file descriptor* to use for this accepted connection and -1 on error

Example

```
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#define MYPORT 3490 // the port users will be connecting to
#define BACKLOG 10 // pending connections queue will hold
main(){
    int sockfd, new_fd; // listen on sock_fd, new connection on
    new_fd
    struct sockaddr_in my_addr; // my address information
    struct sockaddr_in their_addr; // connector's address information
    int sin_size;
    sockfd = socket(PF_INET, SOCK_STREAM, 0);
```

Cont...

```
my_addr.sin_family = AF_INET;      // host byte order
```

```
my_addr.sin_port = htons(MYPORT);  // short, network byte  
order
```

```
my_addr.sin_addr.s_addr = INADDR_ANY; // auto-fill with my IP
```

```
memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct
```

```
// don't forget your error checking for these calls:
```

```
bind(sockfd, (struct sockaddr *)&my_addr, sizeof(struct  
sockaddr));
```

```
listen(sockfd, BACKLOG);
```

```
sin_size = sizeof(struct sockaddr_in);
```

```
new_fd = accept(sockfd, (struct sockaddr *)&their_addr,  
&sin_size);
```

send() and recv() - Let's talk!

- The two functions are for communicating over stream sockets or connected datagram sockets.
- `int send(int sockfd, const void *msg, int len, int flags);`
 - `sockfd` is the socket descriptor you want to send data to (returned by `socket()` or got from `accept()`)
 - `msg` is a pointer to the data you want to send
 - `len` is the length of that data in bytes
 - set `flags` to 0 for now
 - `send()` returns the number of bytes actually sent (may be less than the number you told it to send) or -1 on error

send() and recv() - Let's talk!

- `int recv(int sockfd, void *buf, int len, int flags);`
 - `sockfd` is the socket descriptor to read from
 - `buf` is the buffer to read the information into
 - `len` is the maximum length of the buffer
 - set `flags` to 0 for now
 - `recv()` returns the number of bytes actually read into the buffer or -1 on error
 - If `recv()` returns 0, the remote side has closed connection on you

sendto() and recvfrom() - DGRAM style

- `int sendto(int sockfd, const void *msg, int len, int flags, const struct sockaddr *to, int tolen);`
 - *to* is a pointer to a struct `sockaddr` which contains the destination IP and port
 - *tolen* is `sizeof(struct sockaddr)`
- `int recvfrom(int sockfd, void *buf, int len, int flags, struct sockaddr *from, int *fromlen);`
 - *from* is a pointer to a local struct `sockaddr` that will be filled with IP address and port of the originating machine
 - *fromlen* will contain length of address stored in *from*

Sending message on socket

- ssize_t **sendto**(int sockfd, const void *buf, size_t len, int flags, const struct sockaddr *dest_addr, socklen_t addrlen)

The above call is used to send a message on the socket

- **Arguments :**
 - sockfd – File descriptor of socket
 - buf – Application buffer containing the data to be sent
 - len – Size of buf application buffer
 - flags – Bitwise OR of flags to modify socket behaviour
 - dest_addr – Structure containing address of destination
 - addrlen – Size of dest_addr structure

Receiving message from socket

- `ssize_t recvfrom(int sockfd, void *buf, size_t len, int flags, struct sockaddr *src_addr, socklen_t *addrlen)`

Above call is used to receive message from the socket.

- **Arguments :**
 - `sockfd` – File descriptor of socket
 - `buf` – Application buffer in which to receive data
 - `len` – Size of buf application buffer
 - `flags` – Bitwise OR of flags to modify socket behaviour
 - `src_addr` – Structure containing source address is returned
 - `addrlen` – Variable in which size of `src_addr` structure is returned

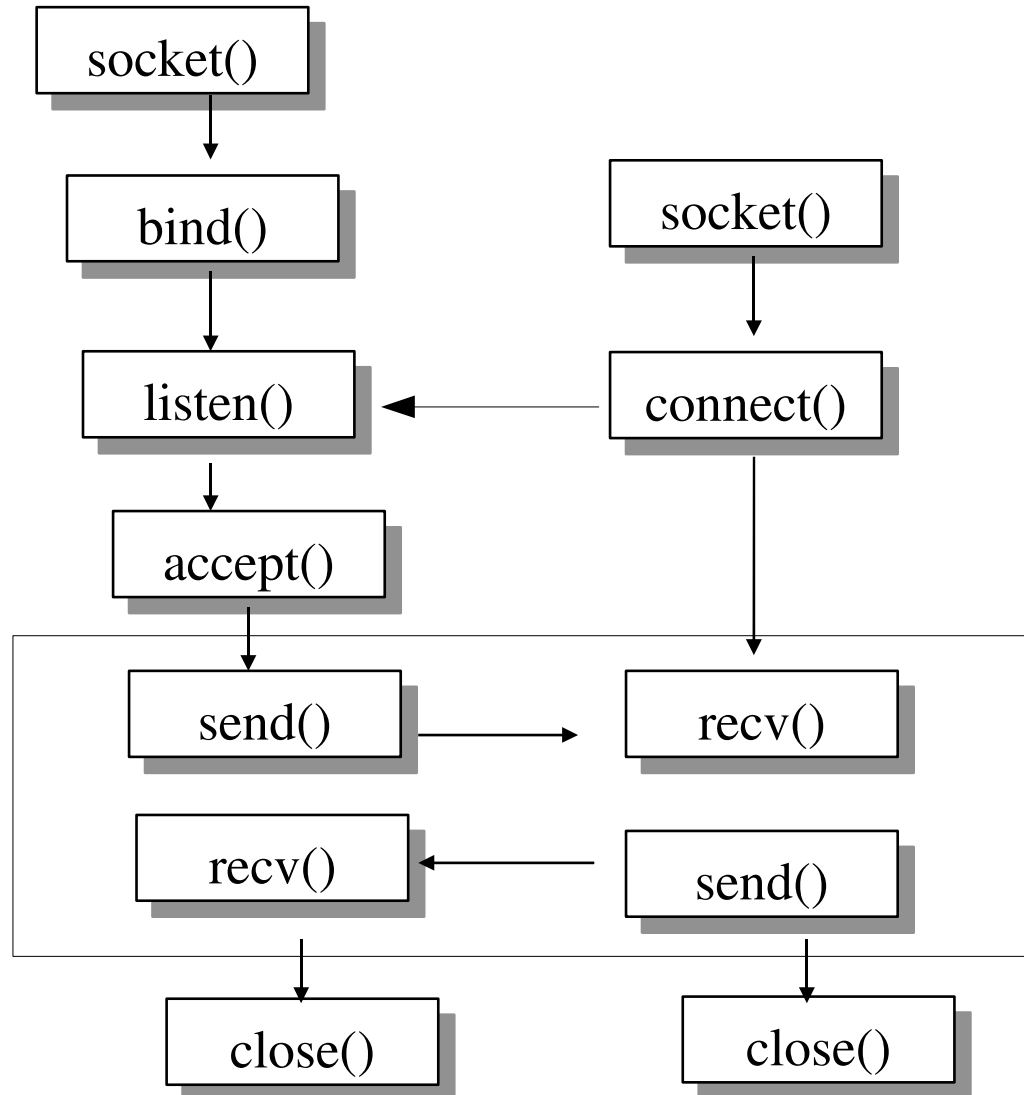
close() - Bye Bye!

- `int close(int sockfd);`
 - Closes connection corresponding to the socket descriptor and frees the socket descriptor
 - Will prevent any more sends and recvs

Connection Oriented Protocol

Server

Client



UDP Client Server Implementation

The entire process can be broken down into following steps :

UDP Server :

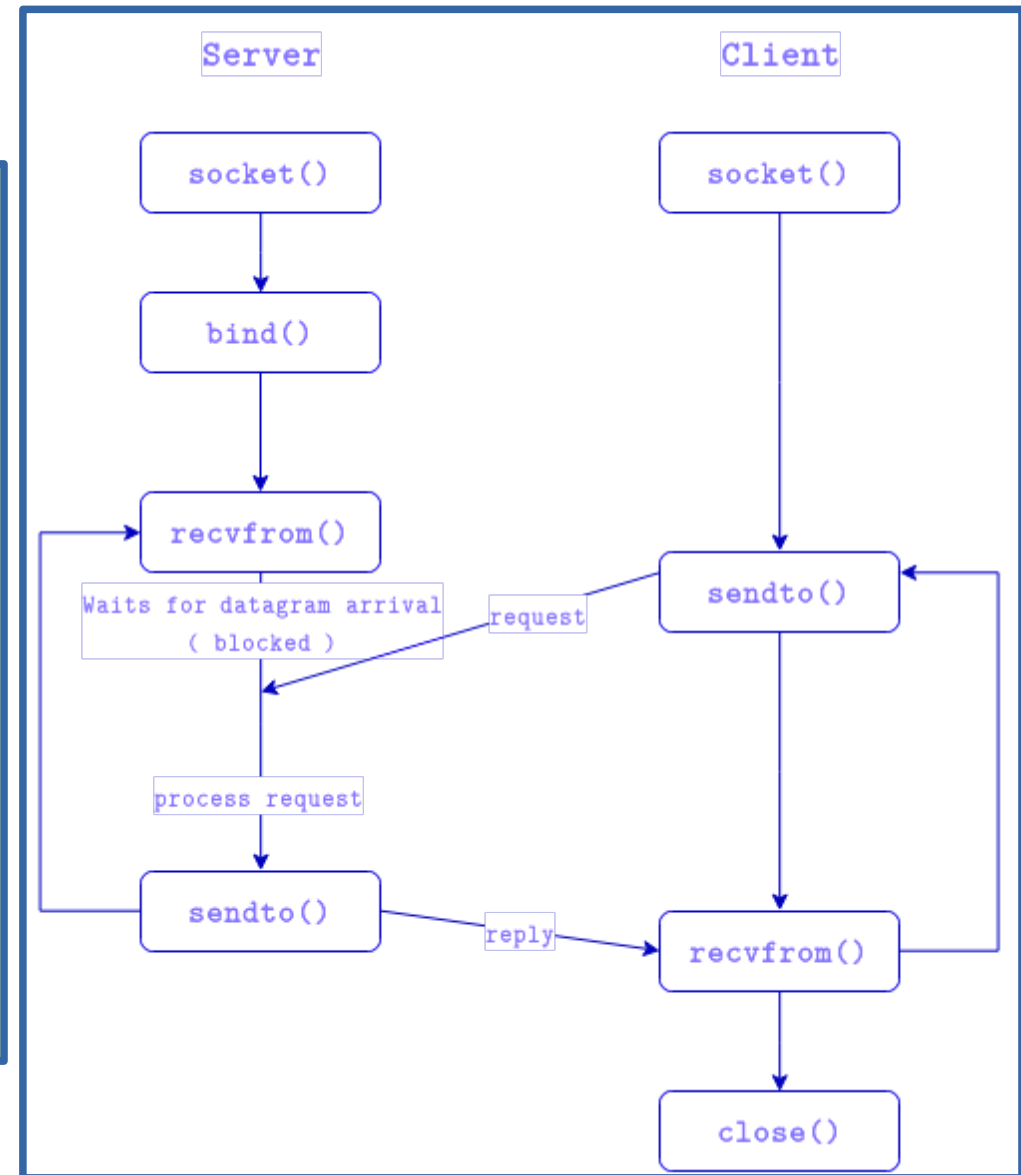
- 1) Create UDP socket.
- 2) Bind the socket to server address.
- 3) Wait until datagram packet arrives from client.
- 4) Process the datagram packet and send a reply to client.
- 5) Go back to Step 3.

UDP Client :

- 1) Create UDP socket.
- 2) Send message to server.
- 3) Wait until response from server is received.
- 4) Process reply and go back to step 2, if necessary.
- 5) Close socket descriptor and exit.

UDP Client-Server Implementation Flow

- In UDP, the client does not form a connection with the server like in TCP and instead just sends a datagram.
- Similarly, the server need not accept a connection and just waits for datagrams to arrive.
- Datagrams upon arrival contain the address of sender which the server uses to send data to the correct client.



Miscellaneous Routines

- `int getpeername(int sockfd, struct sockaddr *addr, int *addrlen);`
 - Will tell who is at the other end of a connected stream socket and store that info in *addr*
- `int gethostname(char *hostname, size_t size);`
 - Will get the name of the computer your program is running on and store that info in *hostname*

Miscellaneous Routines

- `struct hostent *gethostbyname(const char *name);`
 `struct hostent {`
 `char *h_name; //official name of host`
 `char **h_aliases; //alternate names for the host`
 `int h_addrtype; //usually AF_NET`
 `int h_length; //length of the address in bytes`
 `char **h_addr_list; //array of network addresses for the host`
 `}`