

# Threads

a new abstraction of program execution



# Example

```
while (1) {  
    int sock = accept();  
    if (0 == fork()) {  
        handle_request();  
        close(sock);  
        exit(0);  
    }  
}
```

- A web server
  - A process listens for requests
  - When a new request comes in, create a new child process to handle this request
  - Multiple requests can be handled at the same time by different child processes

# Example Web Server

```
while (1) {  
    int sock = accept();  
    if (0 == fork()) {  
        handle_request();  
        close(sock);  
        exit(0);  
    }  
}
```

This implementation is very **inefficient**.

Why?

# Web Server using fork() is inefficient

- **Time**

- creating a **new process** for each incoming request
- **context switching** between processes

- **Space**

- each child process has its **own copy** of the address space, but they are almost the **same**

- **Inter-process communication (IPC)**

- the child process may generate a response which will be used by the parent process
- **Extra overhead** to make this happen (use a file/pipe/socket/shared memory)

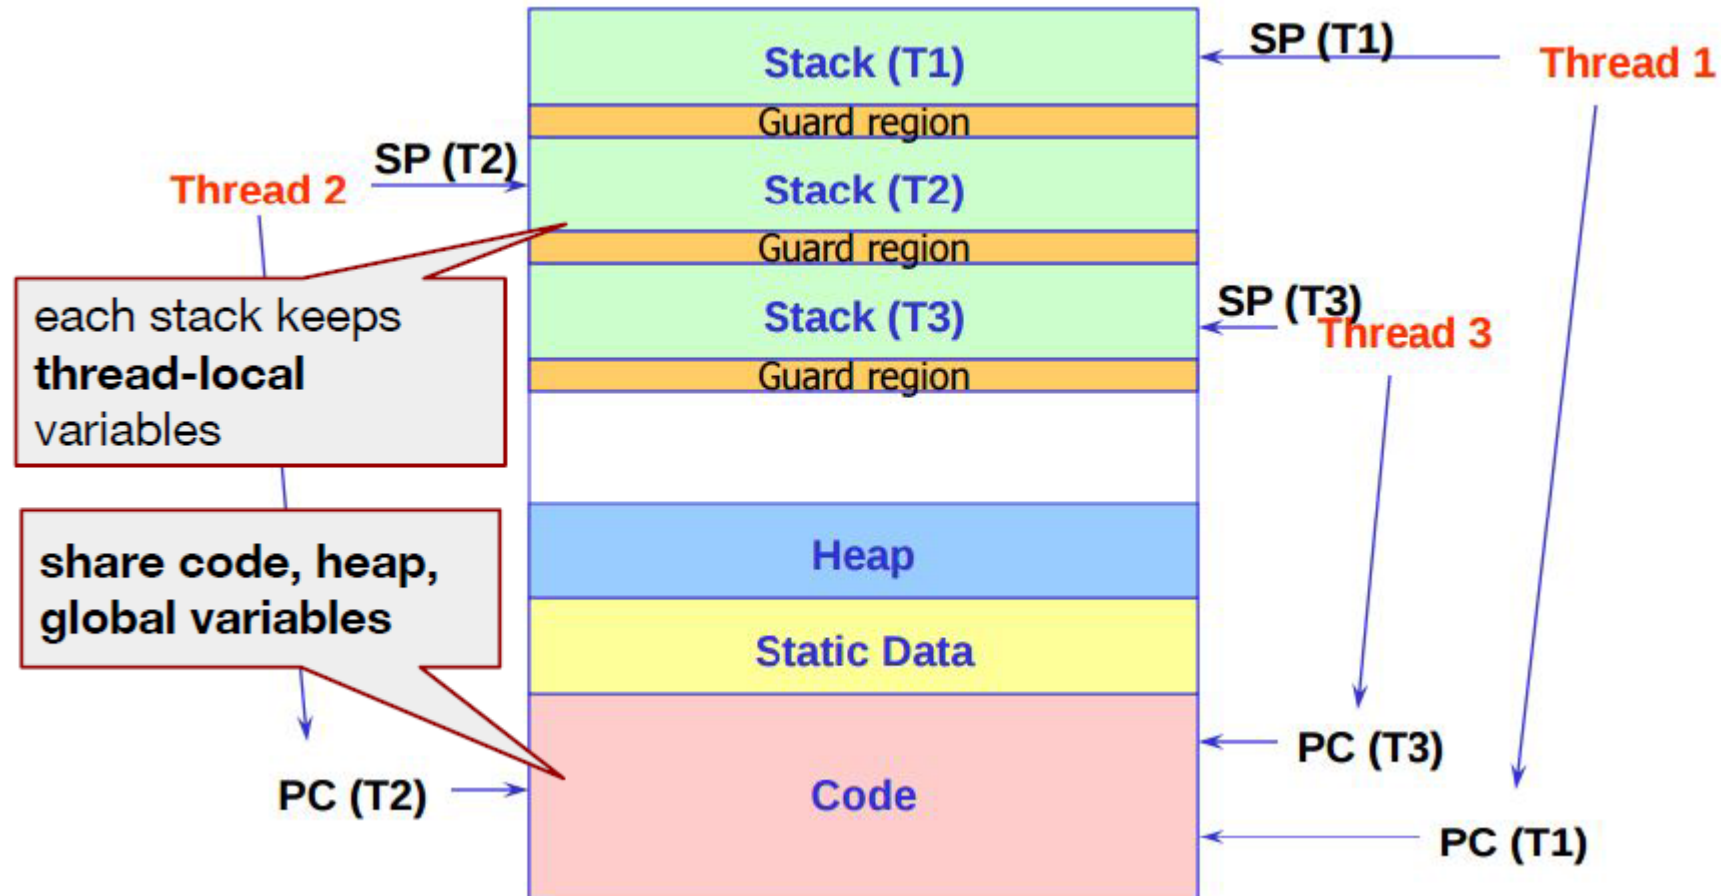
# What is a Thread?

- A normal process is a running program with a **single control flow**, i.e., a single PC.
- A **multi-threaded program** has **multiple control flows**, i.e., multiple PCs.
  - One thread is a single **control flow** through a program

# Multi-threaded Process Address Space

- Multiple threads share the **same address space**
  - no need to copy address space when creating thread
  - no need to switch address space on context switch
- There are **multiple stacks** in the address space, one for each thread.

# Multi-threaded Process Address Space



# Web Server Using Threads

- Different threads share and modify the global variable

```
int global_counter = 0;
web_server() {
    while (1) {
        int sock = accept();
        thread_create(handle_request, sock);
    }
}
```



# The counter thing doesn't work using fork(). Why?

```
int global_counter = 0;
while (1) {
    int sock = accept();
    if (0 == fork()) {
        handle_request();
        ++global_counter;
        close(sock);
        exit(0);
    }
}
```

- counter is always 0 in the parent process, since changes only happen in children
- If you really want to make this work using fork(), need inter-process communication

# Summary of threads

- Lighter weight
  - faster to create and destroy
  - faster context switch
- Sharing
  - threads can solve a single problem concurrently and can easily share code, heap and global variables

# Thread programming

- Linux pthreads using kernel-level threads
- Thread prototypes: `#include <pthread.h>`  
> 60 functions
- To compile:

```
gcc -lpthread prog.c -o prog
```

OR

```
gcc -pthread prog.c -o prog
```

# Basic operations on threads

- **Creating and starting a thread**

- like an asynchronous procedure call

- **Joining with a thread**

- blocks the calling thread until a target thread completes
  - a typical operation when a thread needs to use the return value of the target-thread's starting procedure

# Thread Creation

```
int pthread_create (pthread_t *tid,  
                   pthread_attr_t *attr,  
                   void *(*func)(void *),  
                   void *arg);
```

- **func** : input parameter, the name of the user-defined function to be executed by the newly created thread.
  - returns a void pointer, single void pointer argument
- **arg** : input parameter, the argument to be passed to the function **func**

# Thread Creation

```
int pthread_create (pthread_t *tid,  
                    pthread_attr_t *attr,  
                    void *(*func)(void *),  
                    void *arg);
```

- Why void \* as argument type?
- The argument type should:
  - Work for all types and sizes of arguments
  - Even work for different numbers of arguments

# Thread Creation

```
int pthread_create (pthread_t *tid,  
                   pthread_attr_t *attr,  
                   void *(*func)(void *),  
                   void *arg);
```

- **tid** : output parameter, pointer to location where the id of the newly created thread will be stored.
- **attr** : input parameter, pointer to structure specifying the thread's attributes (e.g. stack size, etc.)
  - If set to NULL, the system defaults are used

# Threaded “hello, world”

```
/*
 * hello.c - Pthreads "hello, world" program
 */
#include "csapp.h"

/* thread routine */
void *thread(void *vargp) {
    printf("Hello, world!\n");
    return NULL;
}

int main() {
    pthread_t tid;

    Pthread_create(&tid, NULL, thread, NULL);
    exit(0);
}
```

Wrong!



# Thread coordination

- Threads can only execute as long as their containing process exists
- In the example, what could happen if the thread routine has not reached **printf** with the main thread reaching the end of `main()`?
- We need to use **join()** to make the main thread wait for the other thread to finish

# pthread\_join()

```
int pthread_join(pthread_t thread,  
                 void **value_ptr );
```

- **thread**: input parameter, id of the thread to wait on
- **value\_ptr**: output parameter, value given to **pthread\_exit()** by the terminating thread (which happens to always be a **void \***)
  - Type is void \*\* because type of location to be updated is void \*

# Threaded “hello, world”

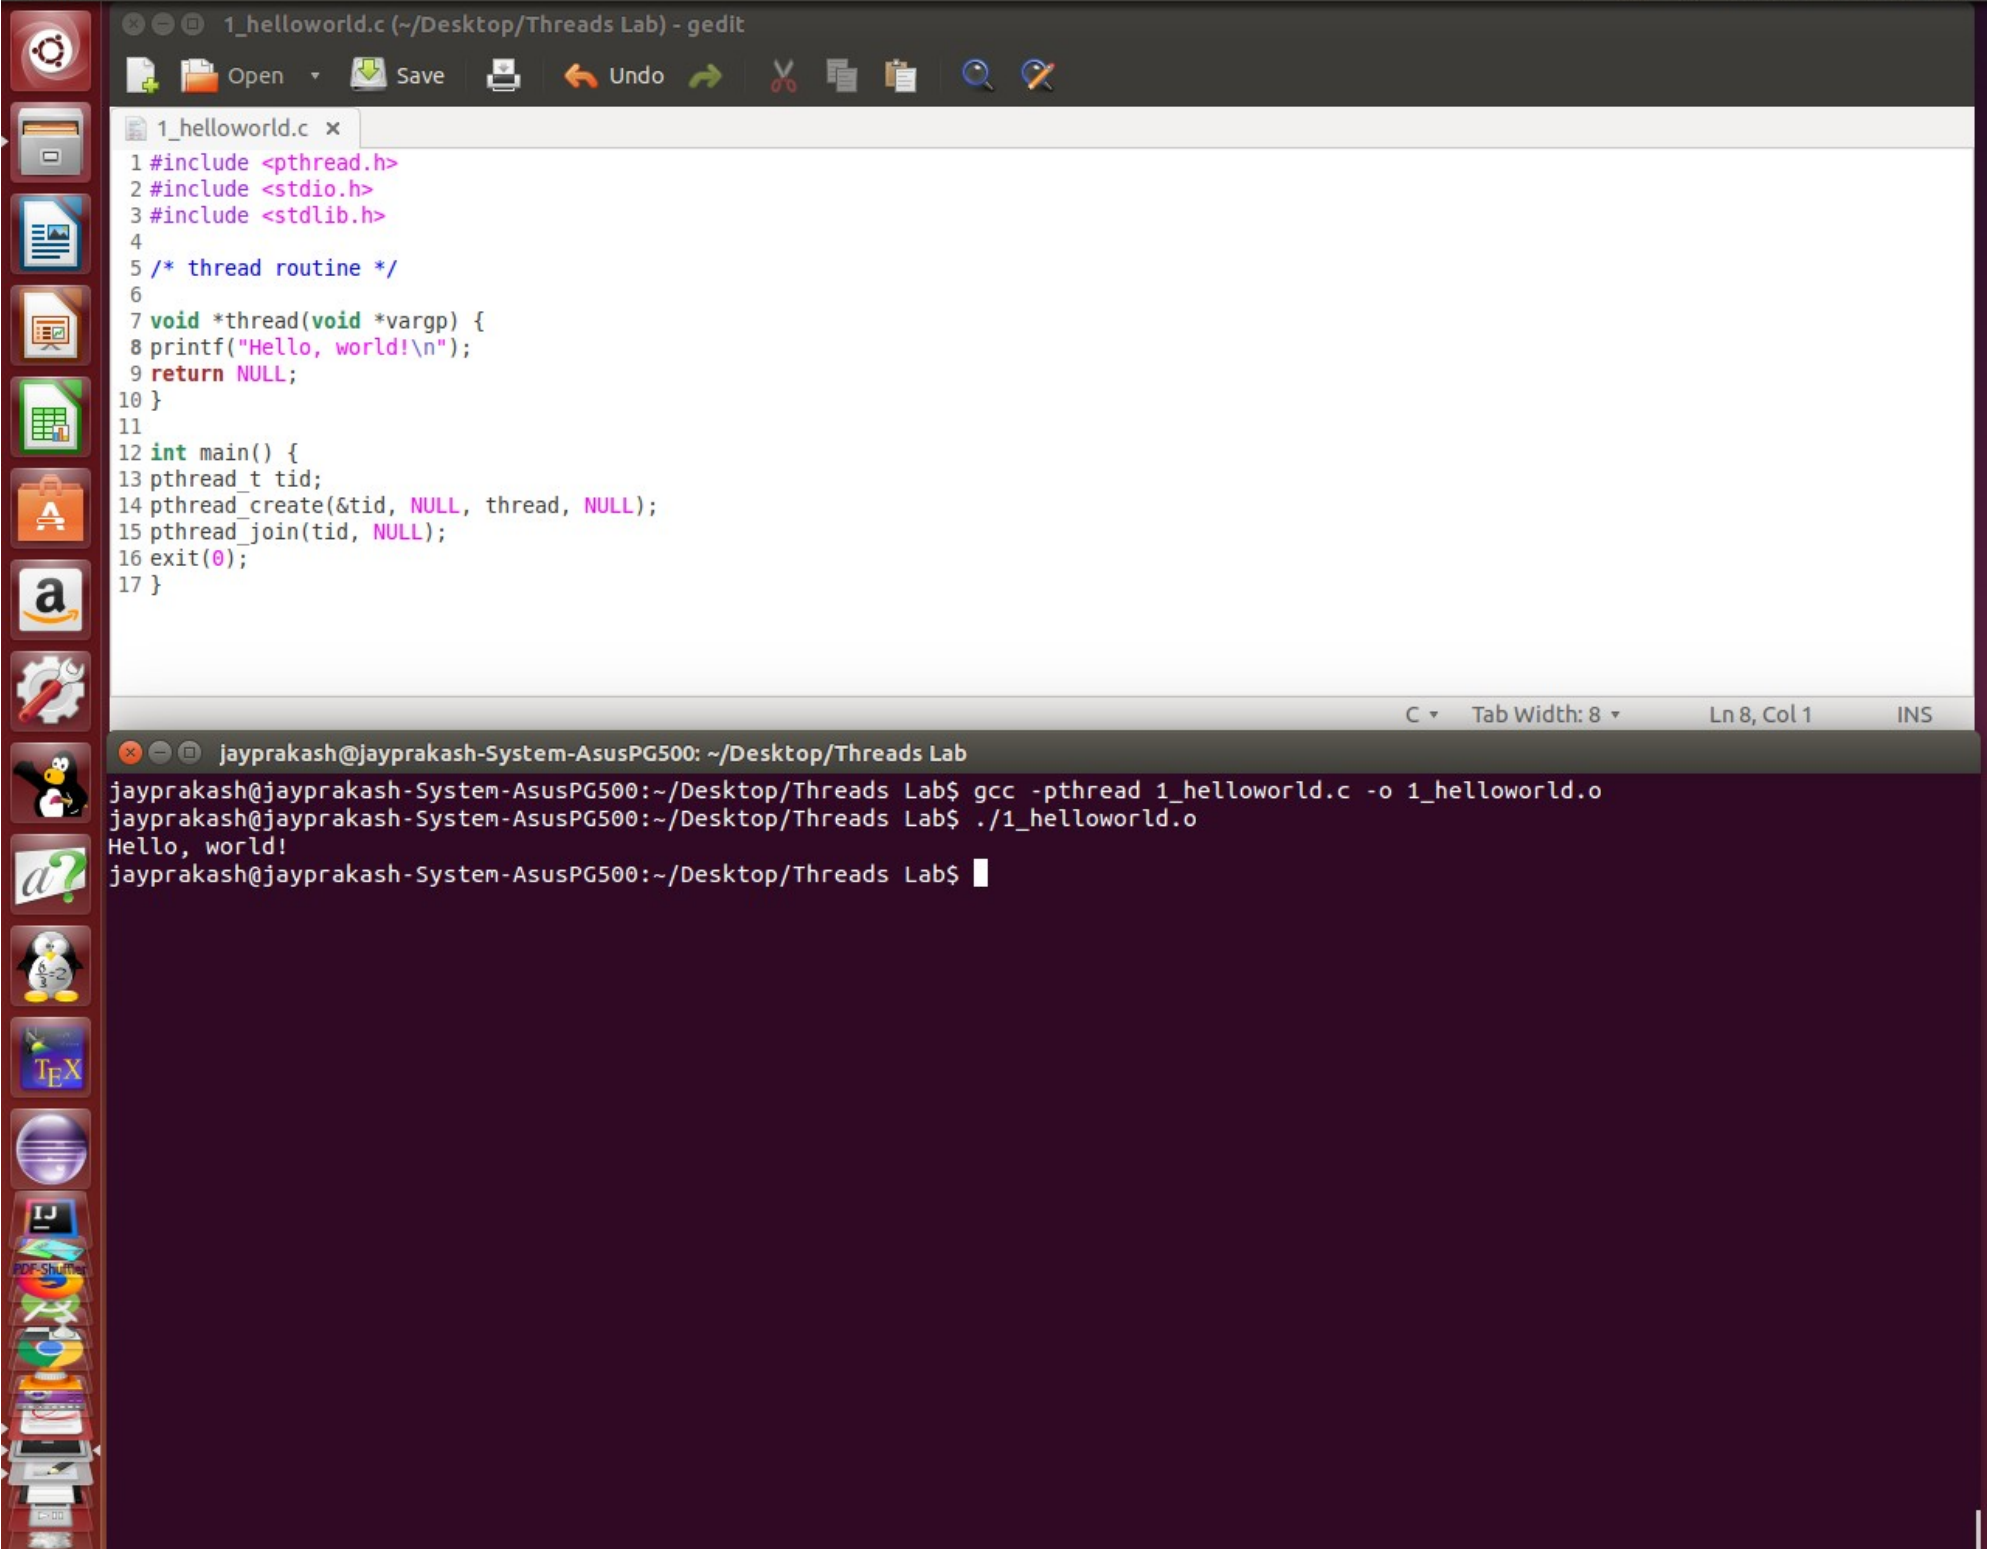
```
/*
 * hello.c - Pthreads "hello, world" program
 */
#include "csapp.h"

/* thread routine */
void *thread(void *vargp) {
    printf("Hello, world!\n");
    return NULL;
}

int main() {
    pthread_t tid;

    Pthread_create(&tid, NULL, thread, NULL);
    Pthread_join(tid, NULL);
    exit(0);
}
```

Good!



# Execution of threaded “hello, world”

