

# Introduction to Mongo DB



pm\_jat @ daiict



# Document Databases - Recap

- Based on Key-Value strategy, where Value is a Document
- A collection of documents is referred as “Collection”
- A document database is collection of “document **collections**”
- Here is RDB and Document DB correspondences

RDB	Doc DB
Database/Schema	Database
Table	“SET” of similar Documents
Row/Tuple	“ <b>Document</b> ”
Row ID	_id



# Mongo DB “Data Models”

- Document as a basic unit
- “Collection” – collection of related documents
- We can have a collection level “schema”
- MongoDB database design uses following strategies
  - Collections with “embedded documents” – less normalized
  - Separated document Collections – more normalized
  - Typically mongo DB uses combined strategy.
- So how does mongo db look like?
  - **“MongoDB Database is collection of document collections”** vis-à-vis “relational database is collection of tables”



# Why do we have “Schema” for databases?

- Schema information helps DBMS in **ensuring Database Integrity Constraints**, that is
  - data are valid in terms of types, cardinality, and other existential constraints like not null, unique, referential integrity etc.
- Schema information acts as assertion that can be used for
  - Efficient storage and access paths (File organization and indexes).
  - Query Optimization and efficient execution of queries



# Motivation for Schema-less/Flexible Schema

- Why Schema-less/Flexible Schema is becoming important?
- In Big Data era: figuring out schema in advance is becoming impractical
- More and more column may get added with the time
- Some columns may be becoming irrelevant with the time
- Creating databases in “fixed schema” approach have problems of
  - Too many columns, and every row having values only for few of them.
  - Database becomes sparse with lots of null values



# Problems with schema-less-ness

- Totally schema-less can be disastrous?
  - Two data objects having two names quantity attribute
    - qty for quantity; how do we reconcile this?
  - “3” and “three” as value for an attributes
- DBMS can not do any kind of validation for such anomalies in absence of schema?
- So, we require having a tradeoff between strict schema and schema less.
- Sometime schema is “**Implicit**” only. That is DBMS does not know anything about schema. Where as “application programmer” can assume some schema and access accordingly.
- Again, DBMS can not perform any kind of validation for implicit schema.
- Also DBMS performance of database operation gets affected in absence of schema.



# MongoDB and Schema-less-ness

- MongoDB's collections, by default, does not require its documents to have the same schema, i.e. documents in a collection
  - Field names can be different
  - data type for a field can differ across documents within a collection
- This means each document in a collection can have its own schema?
- However we can specify some partial schema at collection level? i.e. required fields for every document, their “global data type”, and so on
- “Schema of a collection” (or document) can be altered later also?



# MongoDB and Schema

- We can define certain (Schema) rules at Collection level.
- Mongo DB calls them “Document Validation Rules”.
  - This can be read as Rule that every document in a collection should comply.
- Document Validation rules can be specified at Collection Creation time or can be specified later.
- Mongo DB provides “**JSON Schema**” as the means of defining schema validation rules.
- Here is an JSON Schema example from Mongo Doc Pages

Source: <https://docs.mongodb.com/manual/core/schema-validation/>





# Example: JSON Schema

```
db.createCollection("students", {
  validator: {
    $jsonSchema: {
      bsonType: "object",
      required: [ "name", "year", "major", "address" ],
      properties: {
        name: {
          bsonType: "string",
          description: "must be a string and is required"
        },
        year: {
          bsonType: "int",
          minimum: 2017,
          maximum: 3017,
          description: "must be an integer in [ 2017, 3017 ] and is required"
        },

```



```
major: {
  enum: [ "Math", "English", "Computer Science", "History", null ],
  description: "can only be one of the enum values and is required"
},
gpa: {
  bsonType: [ "double" ],
  description: "must be a double if the field exists"
},
address: {
  bsonType: "object",
  required: [ "city" ],
  properties: {
    street: {
      bsonType: "string",
      description: "must be a string if the field exists"
    },
    city: {
      bsonType: "string",
      "description": "must be a string and is required"
    }
  }
}
```



# Create Collection

```
db.createCollection( <name>,  
  { //options  
    capped: <boolean>,  
    autoIndexId: <boolean>,  
    size: <number>,  
    max: <number>,  
    storageEngine: <document>,  
    validator: <document>,  
    validationLevel: <string>,  
    validationAction: <string>,  
    indexOptionDefaults: <document>,  
    viewOn: <string>, /  
    pipeline: <pipeline>, /  
    collation: <document>, /  
    writeConcern: <document>  
  }  
}
```

Three options related to Schema Validations:

- validator
- validationLevel
- validationAction



# JSON Schema validation options

- **validator** – Used to specify validation rules for the collection. Rules are specified as JSON document containing rules in the form of mongo db query expressions.
- **validationLevel** – specifies “how strictly” MongoDB applies validation rules to existing documents during an update, and
- **validationAction** - specifies whether MongoDB should error and **reject documents** that violate the validation rules or warn about the **violations in the log** but allow invalid documents.



# Validation Level option

- `validationLevel` option values: `off`, `strict`, `moderate`
- It determines on what operations MongoDB applies the validation rules
- If the `validationLevel` is `strict` (the default), MongoDB applies validation rules to all inserts and updates.
- If the `validationLevel` is `moderate`, MongoDB applies validation rules to inserts and to updates to existing documents that already fulfill the validation criteria.
- With the moderate level, updates to existing documents that do not fulfill the validation criteria are not checked for validity.



# Validation Action options

- validationAction option values: error, warn
- It determines how MongoDB handles documents that violate the validation rules:
  - If the validationAction is **error** (the default), MongoDB rejects any insert or update that violates the validation criteria.
  - If the validationAction is **warn**, MongoDB logs any violations but allows the insertion or update to proceed.



# \_id field of Mongo DB documents

- In MongoDB, each document stored in a collection requires a unique \_id field that acts as a primary key.
- If an inserted document omits the \_id field, the MongoDB driver automatically generates an ObjectId for the \_id field.
- The \_id field has the following behavior and constraints:
  - By default, MongoDB creates a unique index on the \_id field during the creation of a collection.
  - The \_id field is always the first field in the documents.
  - The \_id field may contain values of any BSON data type, other than an array.



# BSON data types

- **BSON** is a binary serialization format used to store documents and make remote procedure calls in MongoDB.
- The BSON specification is available at <http://bsonspec.org/>
- These are the data types that can be used for specifying schema and field values





# Mongo DB

- You may not require creating a collection at all.



# Creating Mongo DB Collection

- MongoDB may not require explicit creation of a empty collection.
- First use makes one (implicit creation); for example following will create one: `booksDB.books.insertOne(...)`
- Or can be explicitly created using `booksDB.createCollection(...)`



# Mongo DB Operations

- Mongo DB is very compressive No SQL systems, and it provides very rich set of “Transactional” and “Analytical operations”.
- Typically what operations, a database should have
  - CRUD
  - Transactional Support for concurrent access – ACID?
  - More here: dealing with “shards” and “replicas”



# Mongo DB Operations

- Mongo DB is very compressive No SQL systems, and it provides very rich set of “Transactional” and “Analytical operations”.
- Write operations
  - INSERT, UPDATE, DELETE
  - All write operations in MongoDB are “**atomic**” on the level of a single document.
- Read operation
  - Find
  - JOIN
  - Aggregate and Aggregation Pipeline (we can create customizing sequence of analytical tasks)
  - map-reduce (can also specify map and reduce codes for customized processing) however aggregation pipeline s supposed to be doing more efficient job)



# Read/Write Concerns

- Read/Write Concerns are specified with every read write operations
  - Individually for each operation or Setting for a “Transaction”
- This is the mechanism of setting “Consistency Level”
- For write operations, we typically specify do following in r/w concerns-
  - For **write**, we specify, **how many replicas** a write should be performed before returning success
  - For **read**, we specify, **which replica** do we read from?



# INSERT operation

- Provide following commands
  - `db.collection.insertOne //one`
  - `db.collection.insertMany //many`
  - `db.collection.insert //one or many`
- This **inserts specified documents if not already there in the collection**
- It also creates the collection, if collection not already there
- Value for **\_id field is automatically supplied** if not specified
- All write operations in MongoDB are atomic on the level of a single document.
- We can also specify the “**write concerns**”, i.e. the level of acknowledgement?



# INSERT Examples

```
db.inventory.insertOne(  
  { item: "canvas", qty: 100, tags: ["cotton"], size: { h: 28, w: 35.5, uom: "cm" } }  
)
```

```
db.inventory.insertMany([  
  { item: "journal", qty: 25, size: { h: 14, w: 21, uom: "cm" }, status: "A" },  
  { item: "notebook", qty: 50, size: { h: 8.5, w: 11, uom: "in" }, status: "A" },  
  { item: "paper", qty: 100, size: { h: 8.5, w: 11, uom: "in" }, status: "D" },  
  { item: "planner", qty: 75, size: { h: 22.85, w: 30, uom: "cm" }, status: "D" },  
  { item: "postcard", qty: 45, size: { h: 10, w: 15.25, uom: "cm" }, status: "A" }  
]);
```

Source: <https://docs.mongodb.com/manual/tutorial/insert-documents/>



# Insert – more methods

- The following methods can also add new documents to a collection:
  - `db.collection.update()` when used with the `upsert: true` option.
  - `db.collection.updateOne()` when used with the `upsert: true` option.
  - `db.collection.updateMany()` when used with the `upsert: true` option.
  - `db.collection.findAndModify()` when used with the `upsert: true` option.
  - `db.collection.findOneAndUpdate()` when used with the `upsert: true` option.
  - `db.collection.findOneAndReplace()` when used with the `upsert: true` option.
  - `db.collection.save()`.
  - `db.collection.bulkWrite()`.





# Query Documents - FIND

- Mongo FIND is very comprehensive like **SQL SELECT**
- It is used as: “**db.collection-name.find(...)**”
- We specify query selection criteria as parameter to find in the form of “json document”
- To get all documents we leave criteria document as empty.  
Example below:

```
db.inventory.find( {} )
```

```
SELECT * FROM inventory
```

Source: <https://docs.mongodb.com/manual/tutorial/query-documents/>



# FIND: specify equality check

- Specified as attribute value pairs. Examples below

```
db.inventory.find( { status: "D" } )
```

```
SELECT * FROM inventory WHERE status = "D"
```

- We can AND by simply using comma, as following:  

```
db.inventory.find( {status: "A", category: 5 } )
```
- For OR, we require little extra work  

```
db.inventory.find(  
    {$or:[{status:"A"}, {status:"D"}] }  
)
```

Source: <https://docs.mongodb.com/manual/tutorial/query-documents/>



# FIND: more examples

- More conditional operators

```
db.inventory.find( { status: "A", qty: { $lt: 30 } } )
```

```
SELECT * FROM inventory WHERE status = "A" AND qty < 30
```

```
db.inventory.find( { $or: [ { status: "A" }, { qty: { $lt: 30 } } ] } )
```

```
SELECT * FROM inventory WHERE status = "A" OR qty < 30
```

Source: <https://docs.mongodb.com/manual/tutorial/query-documents/>



# FIND: more examples

- More conditional operators

```
db.inventory.find( { status: { $in: [ "A", "D" ] } } )  
  
SELECT * FROM inventory WHERE status in ("A", "D")
```

```
db.inventory.find( {  
  status: "A",  
  $or: [ { qty: { $lt: 30 } }, { item: /^p/ } ] //regular expression  
} )  
  
SELECT * FROM inventory WHERE status = "A" AND ( qty < 30 OR item LIKE "p%")
```

Source: <https://docs.mongodb.com/manual/tutorial/query-documents/>



# Embedded/Nested Documents fields in FIND

- Example:  

```
db.inventory.find(  
    { "size.h": { $lt: 15 }, "size.uom": "in",  
      status: "D" }  
)
```
- Predicate expressed here is:  
 $\text{size.h} < 15 \text{ AND size.uom} = \text{"IN"} \text{ AND status} = \text{"D"}$



# Array Field in FIND

- Consider following collection with Array Field `tags`, `dim_cm`

```
db.inventory.insertMany([
  { item: "journal", qty: 25, tags: ["blank", "red"], dim_cm: [ 14, 21 ] },
  { item: "notebook", qty: 50, tags: ["red", "blank"], dim_cm: [ 14, 21 ] },
  { item: "paper", qty: 100, tags: ["red", "blank", "plain"], dim_cm: [ 14, 21 ] },
  { item: "planner", qty: 75, tags: ["blank", "red"], dim_cm: [ 22.85, 30 ] },
  { item: "postcard", qty: 45, tags: ["blue"], dim_cm: [ 10, 15.25 ] }
]);
```

- Query “`db.inventory.find( { tags: ["red", "blank"] } )`” shall look for document that have tags as exactly this two element array, i.e. `["red", "blank"]`



# Array Field in FIND

- Consider following collection with Array Field `tags`, `dim_cm`

```
db.inventory.insertMany([
  { item: "journal", qty: 25, tags: ["blank", "red"], dim_cm: [ 14, 21 ] },
  { item: "notebook", qty: 50, tags: ["red", "blank"], dim_cm: [ 14, 21 ] },
  { item: "paper", qty: 100, tags: ["red", "blank", "plain"], dim_cm: [ 14, 21 ] },
  { item: "planner", qty: 75, tags: ["blank", "red"], dim_cm: [ 22.85, 30 ] },
  { item: "postcard", qty: 45, tags: ["blue"], dim_cm: [ 10, 15.25 ] }
]);
```

- Query “`db.inventory.find( { tags: { $all: ["red", "blank"] } } )`” shall look for document that have this array (i.e. `["red", "blank"]`) as subset of elements in tags attributes.



# Array Field in FIND

```
db.inventory.insertMany([
  { item: "journal", qty: 25, tags: ["blank", "red"], dim_cm: [ 14, 21 ] },
  { item: "notebook", qty: 50, tags: ["red", "blank"], dim_cm: [ 14, 21 ] },
  { item: "paper", qty: 100, tags: ["red", "blank", "plain"], dim_cm: [ 14, 21 ] },
  { item: "planner", qty: 75, tags: ["blank", "red"], dim_cm: [ 22.85, 30 ] },
  { item: "postcard", qty: 45, tags: ["blue"], dim_cm: [ 10, 15.25 ] }
]);
```

- Query “`db.inventory.find( { tags: "red" } )`” returns the documents that have “red” as an element in tags attribute
- Query “`db.inventory.find( { dim_cm: { $gt: 25 } } )`” returns the documents that have at-least one element > 25 in dim\_cm array attribute.





# Find returns “Cursor”

- And you can iterate through as following.  
Note: selection criteria here is “type=2”

```
var myCursor = db.users.find( { type: 2 } );  
  
while (myCursor.hasNext()) {  
    print(tojson(myCursor.next()));  
}
```



# Find: Null or Missing Fields

- If attribute (field) item has null value

```
db.inventory.find( { item: null } )
```

- If item type (BSON) is 10 (i.e. Null)

```
db.inventory.find( { item : { $type: 10 } } )
```

- If attribute item is non existent in a document

```
db.inventory.find( { item : { $exists: false } } )
```



# “read-concern” with FIND

- We can also specify “**read concern**” for **choosing isolation level** while reading from replica sets and replica set shards!
- Typically specifies how many nodes, read should consult before returning a value because, replicate sets can be in middle of update, and all replicas are not consistent?
- We shall talk later about this.



# Projection (specifying fields)

- By default all fields are projected
- We can either be field inclusive mode, or exclusive mode
- `Id:0, name:1, qty: 1`, will show name field only (exclude `_id`, and include name)
- `salary:0, email:0`; will exclude these two fields (rest will be projected)
- 0 and 1 can not be mixed in a project expression except with `_id` field



# UPDATE operation

- Either we use `db.collection.update(<filter>, <update>, <options>)`
- Or one of following variations
  - `db.collection.updateOne(<filter>, <update>, <options>)`
  - `db.collection.updateMany(<filter>, <update>, <options>)`
  - `db.collection.replaceOne(<filter>, <update>, <options>)`
- The method uses operators such as **\$set** and **\$unset** for specifying the updates
- `updateOne()` finds the first document that matches the filter and applies the specified update modifications.
- By default `update()` also updates single document. Has various options to do the same done by all other variations.

Source: <https://docs.mongodb.com/manual/tutorial/update-documents/>



# Update Example

```
db.books.update(  
  { _id: 1 },  
  {  
    $inc: { stock: 5 },  
    $set: {  
      item: "ABC123",  
      "info.publisher": "2222",  
      tags: [ "software" ],  
      "ratings.1": { by: "xyz", rating: 3 }  
    }  
  }  
)
```

```
UPDATE books  
SET  
    stock = stock + 5  
    item = "ABC123"  
    publisher = 2222  
    pages = 430  
    tags = "software"  
    rating_authors = "ijk,xyz"  
    rating_values = "4,3"  
WHERE _id = 1
```

Source: <https://docs.mongodb.com/manual/reference/method/db.collection.update>



# option with Update operation

- **multi**: boolean. by default updates one.
- **upsert**: boolean. Value true creates a new document when no document matches the query criteria.
- **writeConcern**: A document expressing the write concern. Default write concern “w: 1”.
- **hint**: a document or string that specifies the index to use to support the query predicate.



# Update – replace entire document

```
db.books.update(  
  { _id: 2 },  
  {  
    item: "XYZ123",  
    stock: 10,  
    info: { publisher: "2255", pages: 150 },  
    tags: [ "baking", "cooking" ]  
  }  
)
```

```
DELETE from books WHERE _id = 2  
INSERT INTO books  
  (_id,  
   item,  
   stock,  
   publisher,  
   pages,  
   tags)  
VALUES (2,  
        "xyz123",  
        10,  
        "2255",  
        150,  
        "baking,cooking")
```

Source: <https://docs.mongodb.com/manual/reference/method/db.collection.update>





# Update Example: document before and after update

```
"_id" : 1,  
"item" : "TBD",  
"stock" : 0,  
"info" : { "publisher" : "1111", "pages" : 430 },  
"tags" : [ "technology", "computer" ],  
"ratings" : [ { "by" : "ijk", "rating" : 4 }, { "by" : "lmn", "rating" : 5 } ]  
"reorder" : false  
,
```

```
"_id" : 1,  
"item" : "ABC123",  
"stock" : 5,  
"info" : { "publisher" : "2222", "pages" : 430 },  
"tags" : [ "software" ],  
"ratings" : [ { "by" : "ijk", "rating" : 4 }, { "by" : "xyz", "rating" : 3 } ],  
"reorder" : false
```

Source: <https://docs.mongodb.com/manual/reference/method/db.collection.update>



# Update Example with “upsert”

```
db.books.update(  
  { item: "ZZZ135" },    // Query parameter  
  {                      // Replacement document  
    item: "ZZZ135",  
    stock: 5,  
    tags: [ "database" ]  
  },  
  { upsert: true }      // Options  
)
```

Source: <https://docs.mongodb.com/manual/reference/method/db.collection.update>



# Use of Index while updating!

- `hint` option specifies the index (available since ver 4.2)

```
db.members.update(  
  { points: { $lte: 20 }, status: "P" },      // Query parameter  
  { $set: { misc1: "Need to activate" } },    // Update document  
  { multi: true, hint: { status: 1 } }        // Options  
)
```

- Creating Index

```
db.members.createIndex( { status: 1 } )  
db.members.createIndex( { points: 1 } )
```

Source: <https://docs.mongodb.com/manual/reference/method/db.collection.update>



# Update: more examples

```
db.inventory.updateOne(  
  { item: "paper" },  
  {  
    $set: { "size.uom": "cm", status: "P" },  
    $currentDate: { lastModified: true }  
  }  
)
```

```
db.inventory.updateMany(  
  { "qty": { $lt: 50 } },  
  {  
    $set: { "size.uom": "in", status: "P" },  
    $currentDate: { lastModified: true }  
  }  
)
```

```
db.inventory.replaceOne(  
  { item: "paper" },  
  { item: "paper", instock: [ { warehouse: "A", qty: 60 }, { warehouse: "B", qty: 40 } ] }  
)
```



# Delete Documents

- Provide following methods

[db.collection.deleteMany\(\)](#)

[db.collection.deleteOne\(\)](#)

- Examples:

```
db.inventory.deleteMany({}) //delete all
db.inventory.deleteMany({ status : "A" })
                        //delete all with status="A"
db.inventory.deleteOne( { status: "D" } )
                        //delete one with status="A"
```



# References

- [1] Chapter 9, NoSQL Distilled
- [2] <https://docs.mongodb.com/manual/>