

Mongo DB - Indexes



pm_jat @ daiict



A important note!

- There are cursor methods
 - count, max, sort
- There are also \$COUNT, \$MAX, \$SORT in aggregation pipeline.
- Both sound like similar in terms of functionality; however, are different in terms of execution space.
- Cursor methods are “mongo shell” only operations and are not available in API
- That means “cursor methods” are
 - Not distributed (executed on client)
 - Not under the scope of “mongo query optimizer”




Why do we need indexes in Databases?

- Single Reason: **“to make searching most efficient”**
- Availability of indexes also helps in getting the resulted **sorted!**



B+ tree index?






B+ tree “Node”

P_1	K_1	P_2	...	P_{n-1}	K_{n-1}	P_n
-------	-------	-------	-----	-----------	-----------	-------

- Non-leaf node
 - Non leaf nodes contain pointers to its children nodes
 - A pointer P_i in a node points to record with search key $< K_i$ and $\geq K_{i-1}$
 - A search value appearing parent node will also appear in child node.
- Leaf nodes
 - Leaf nodes only store address for data records
 - All leaf nodes combined will be in the order of search key.
 - The pointer P_n (that is last pointer) of a leaf node point to next leaf node, and maintains a linked list of leaf nodes.
 - Leaf node may also be linked doubly; may have a pointer like P_0 referring to previous node.

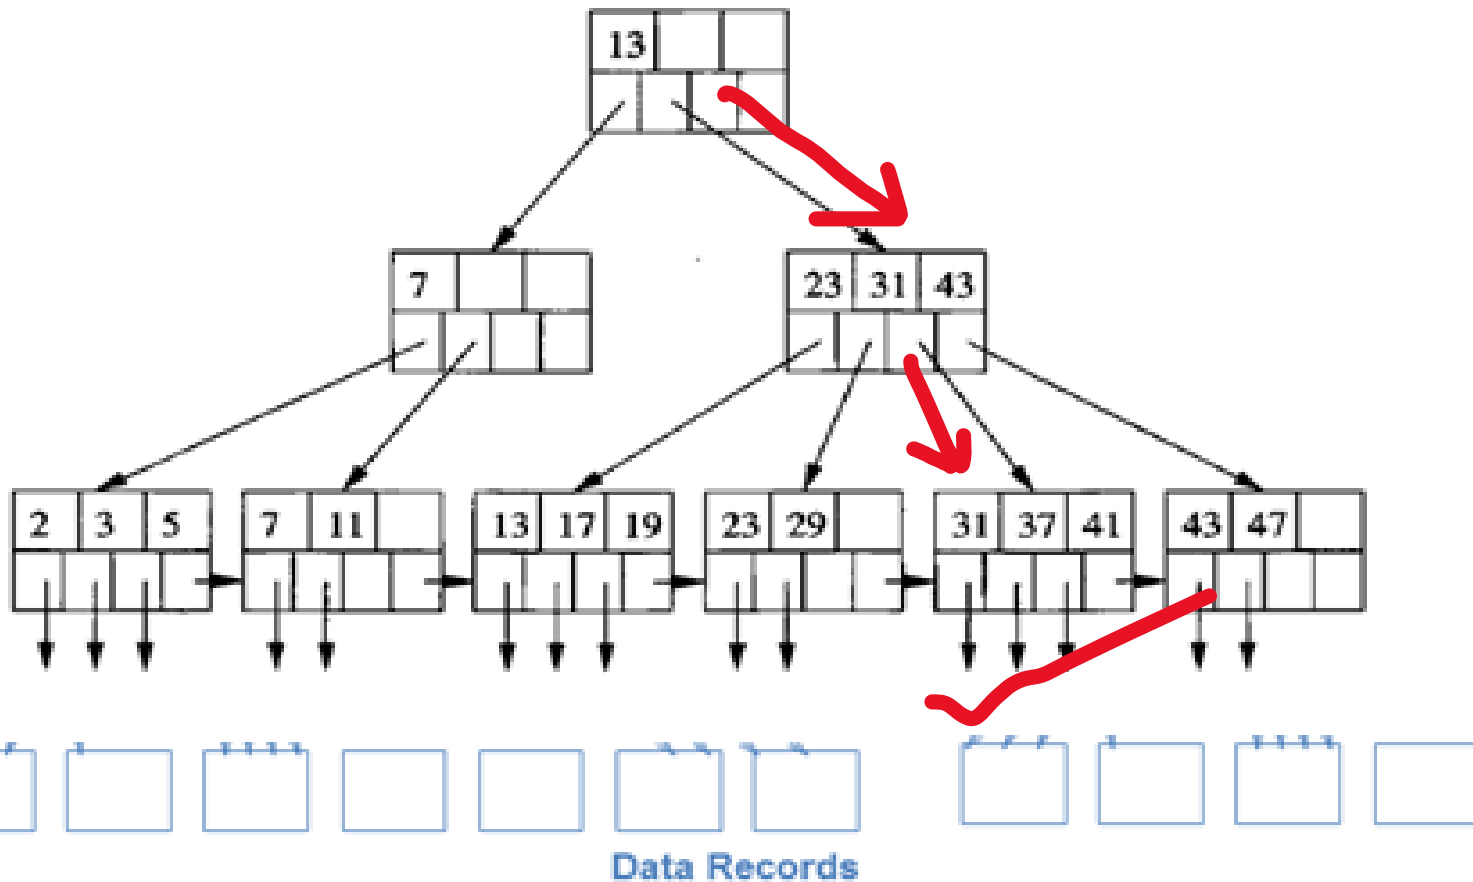
B+ tree index: toy example



Data Records

Let us say this is index for items. Let search key here be item-id.
Can you trace how item id = 31 be searched?

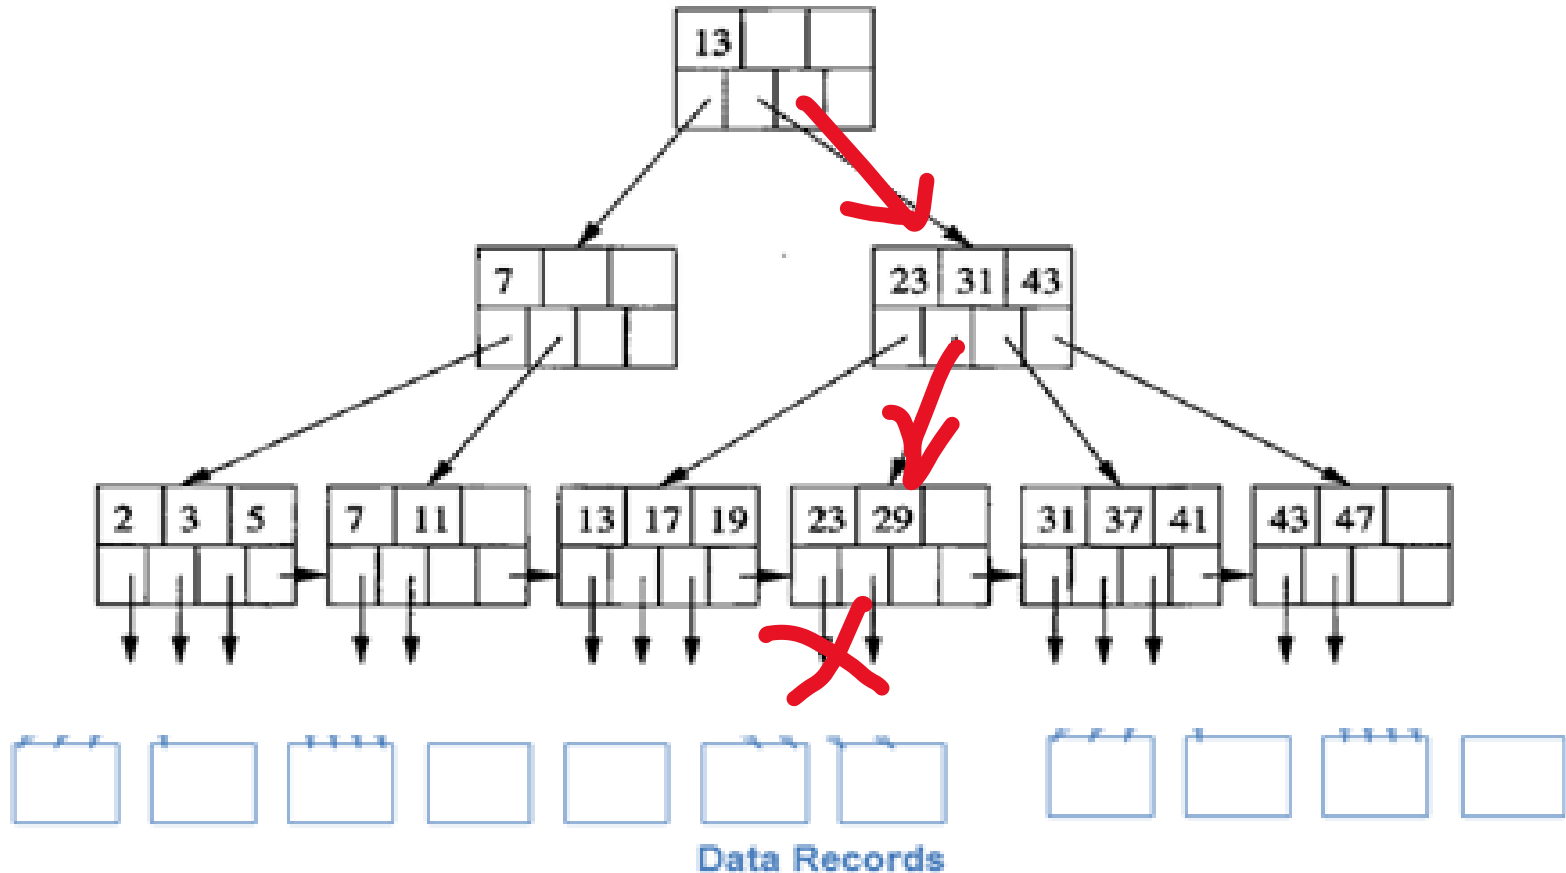
B+ tree index: toy example



Access path for item id = 31




B+ tree index: toy example



Access path for item id = 30?




B+ tree index: toy example



Data Records

Traversal path for item id ≤ 30 ?

B+ tree index: toy example



Data Records

Traversal for $10 \leq x \leq 30$... ascending order

Traversal for $10 \leq x \leq 30$... descending order



Indexes and Query execution

- Hope you know indexes are used in Relational Databases?
- How will indexes affect execution of following query?
 1. `SELECT * FROM PRODUCTS WHERE CAT=5`
 2. `SELECT * FROM PRODUCTS WHERE CAT=5 and PRICE >=2000 and PRICE < 5000`
 3. `SELECT * FROM CUSTOMERS WHERE STATE="ABCD"`
 4. `SELECT * FROM CUSTOMERS WHERE ZIP="122345"`
- Availability of indexes on attributes in predicates makes query execution efficient?



Indexes and Query execution

- For instance: `SELECT * FROM PRODUCTS WHERE CAT=5`
- For getting PRODUCTS of given CAT, we have following options
 - First: Sequentially scan all products and select the one that have CAT=5. BUT this is brute force and not efficient.
 - Two: Sort and Search on CAT
 - Third: TABLE PRODUCTS have index on CAT.
- An index on CAT makes this query executing this query most efficient.
- Interestingly DBMS uses indexes transparently, and we do not have to tell if index is to be used.



Indexes and Query execution

- Query: `SELECT * FROM PRODUCTS WHERE CAT=5 and PRICE >=2000 and PRICE < 5000`
- Best case is when we have composite index on both attributes
- It can still be efficient if there is index on one of attribute (say CAT)
 - Index lookup on indexing field (CAT in this case) and scan for other field
- (Sequential | Sort and Search) when no index



Good and Downside of Indexes

- Good:
 - Helps in improving searches
- Downside:
 - Since DBMS require updating indexes also when data items are getting added and updated.
 - This affects the performance of update operations and hence update operations become slower.
- Therefore, we can not have index on every field and we can not afford having no index at all?



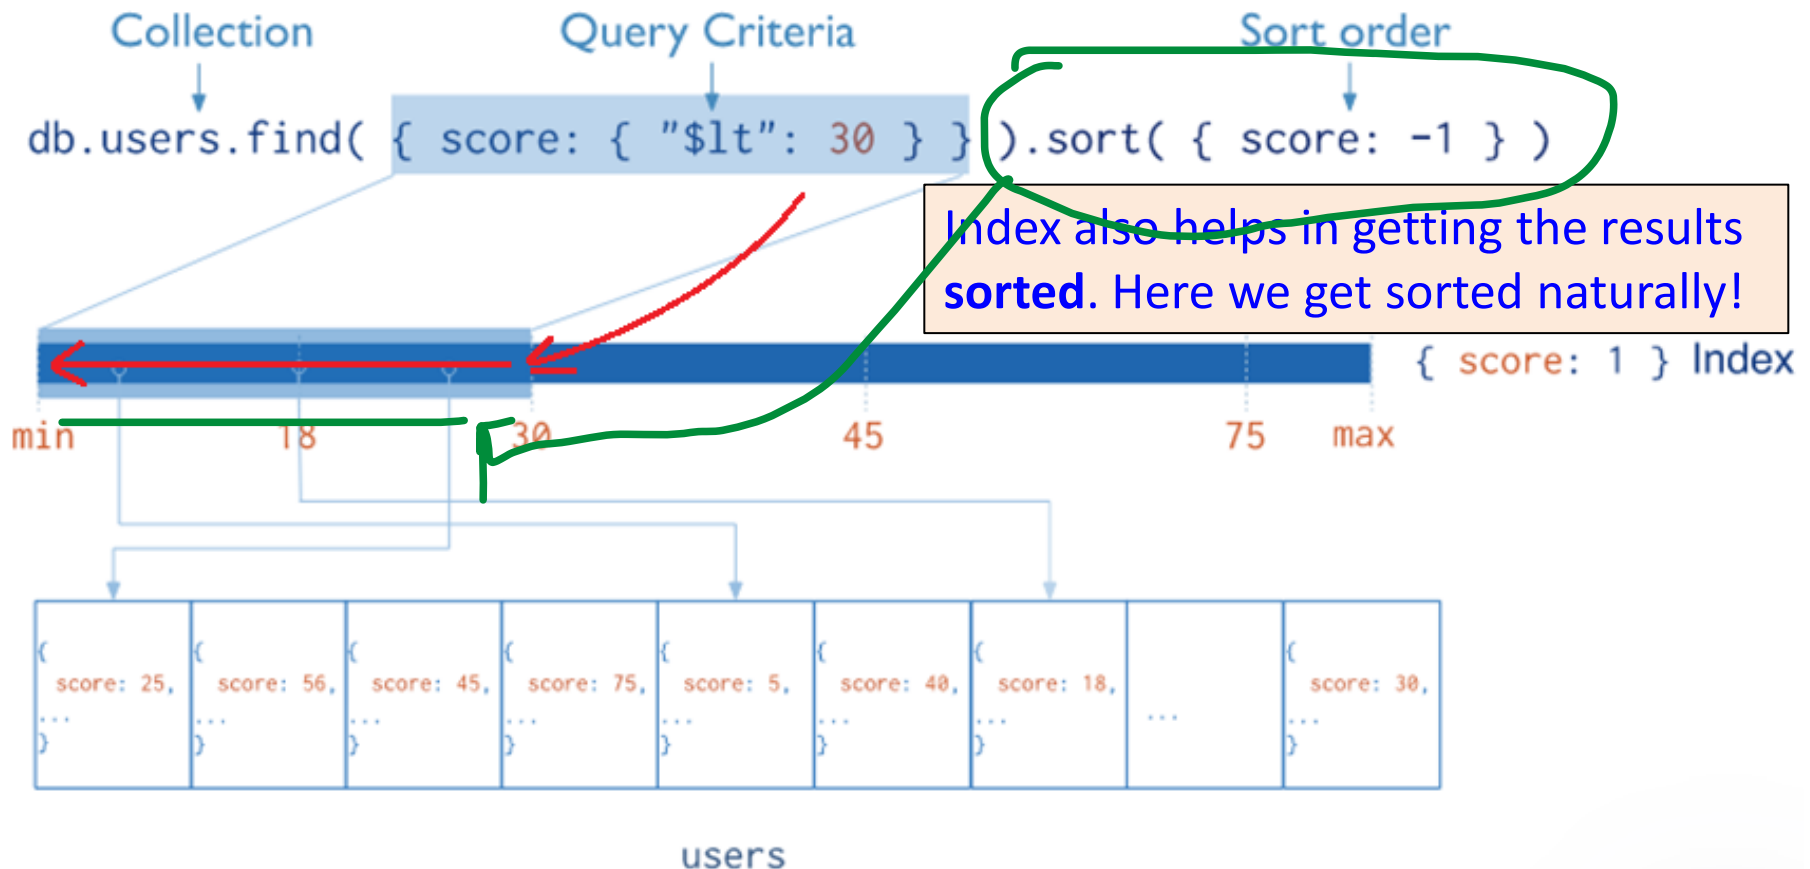
Mongo DB Indexes

- Mongo provides very comprehensive support for Indexes
- Mongo DB indexes are B+ tree based.
 - However, you can also create Hashed Indexes. Hashed Indexes are only good for equality searches
- Mongo DB automatically makes use of indexes wherever indexes are available!



Mongo DB Indexes

- How availability of index on score (ascending) on users will be used in answering following query?



Source: <https://docs.mongodb.com/manual/indexes/>



Default **_id Index**

- MongoDB creates a unique index on the `_id` field during the creation of a collection.
- The `_id` index prevents clients from inserting two documents with the same value for the `_id` field.
- You cannot drop this index on the `_id` field.

Source: <https://docs.mongodb.com/manual/indexes/>



Index Properties

Unique Indexes

- A unique index causes MongoDB to reject duplicate values for the indexed field.

Partial Indexes

- Partial indexes only index the documents in a collection that meet a specified filter expression.
- By indexing a subset of the documents in a collection, partial indexes have lower storage requirements and reduced performance costs for index creation and maintenance.

Source: <https://docs.mongodb.com/manual/indexes/>



Index Properties

TTL Indexes

- TTL (Time To Live) Indexes lives only for specified time.

Source: <https://docs.mongodb.com/manual/indexes/>



Creating Index

- Mongo DB provides method `createIndex`
- Parameters: key and other index specifications, and options

```
db.collection.createIndex( <key and index type specification>, <options> )
```

- Index on zip (ascending) for zips collection:
- Index on state, population for zips collection:

```
db.zips.createIndex( { zip: 1 } )
```

```
db.zips.createIndex(  
  { state: 1, pop: 1 },  
  { name: "stateindex"  
})
```

Source: <https://docs.mongodb.com/manual/indexes/>



Index Types

- MongoDB provides a number of index types to support specific types of data and queries.
- Single Field: On a single field: ascending/descending
- Example:

Index entries (leaf nodes) are sorted on single field
SCORE



{ score: 1 } Index

Source: <https://docs.mongodb.com/manual/indexes/>




Index Types: Compound Index

- Indexes on multiple fields.

Index entries (leaf nodes) are sorted on two field
User, Score

collection



{ userid: 1, score: -1 } Index

Note this index is no more efficient for query

```
db.users.find( { score: { "$lt": 30 } } )
```

Source: <https://docs.mongodb.com/manual/indexes/>



Index Types: Compound Index

- So index on User and Score on Collection Users will be useful for query like
 - find users where userid=101 and score > 50, and
- Not for query like
 - find users with score > 50



“Covered Queries”

- Wishfully all queries are “covered”
- **A covered query is a query that can be satisfied entirely using an index** and does not require looking into documents.
- An index covers a query when all of the following apply:
 - all the fields in the query are part of an index, and
 - all the fields returned in the results are in the same index.
 - no fields in the query are equal to null (i.e. {"field" : null} or {"field" : {\$eq : null}}).

Source: <https://docs.mongodb.com/manual/indexes/>



Some more index types

- **Geospatial Index:** To support efficient queries of geospatial coordinate data
- **Text Indexes:** to support full text search. Index/Searching is done based root words, etc.
- **Hashed Indexes:** To support hash based sharding, MongoDB provides a hashed index type, which indexes the hash of the value of a field. These indexes have a more random distribution of values along their range, but only support equality matches and cannot support range-based queries.

Source: <https://docs.mongodb.com/manual/indexes/>



Dropping indexes

- By field specifications specified in create index, as
`db.zips.dropIndex({ zip: 1 })`
- Or by names if we specified name in create index:
`db.zips.dropIndex("stateindex")`



Full Text Search | Text Search

- Full Text Searching (or just text search) provides the capability to search in natural-language documents
- Keyword based search is quite popular looking into textual descriptions. Consider following two examples
- (1) Searching in Book titles based on input keywords:
 - Suppose we want to search books with following two keywords in title “Mongo+Java”
 - How do we perform? Probably we can use some RegEx but regular expressions are not index friendly!
- (2) Keyword based search in product descriptions?



Full Text Search | Text Search

- The most common type of search is to find all documents containing given query terms and
 - (1) return them in order of their similarity to the query, or
 - (2) Return only top k in terms of similarity
- Notion of query and similarity are very flexible and depend on the specific application.
- The simplest search considers query as a set of words and similarity as the frequency of query words in the document (popular metric is cosine similarity based on Term Frequency and TFIDF)



Text Indexes in Mongo DB

- MongoDB provides text indexes to support basic support for full text searches.
- Text indexes can include any field whose value is a string or an array of string elements.
- Here is an example creating text index on reviews collection on comments fields:
`db.reviews.createIndex({comments:"text"})`



Searching Text index : Toy Example

- Suppose we have collection comments with following documents

```
{ "_id" : "1", "statement" : "MongoDB is the worst." }  
{ "_id" : "2", "statement" : "MongoDB is the best" }  
{ "_id" : "3", "statement" : "DynamoDB is the worst" }  
{ "_id" : "4", "statement" : "DynamoDB is the best" }
```

- And let us create text index as
`db.comments.createIndex({statement: "text"})`
- And let us try search for keywords “MongoDB Best”?



Searching Text index : Toy Example

- Now if we attempt searching:

```
db.comments.find(  
  { $text: { $search: "MongoDB Best" } }  
)
```

- Returns following. Does it surprises?

```
{ "_id" : "4", "statement" : "DynamoDB is the best" }  
{ "_id" : "2", "statement" : "MongoDB is the best" }  
{ "_id" : "1", "statement" : "MongoDB is the worst." }
```



Searching Text index : Toy Example

- Following query results as given below:

```
db.textExample.find({  
  $text: { $search : "MongoDB best" } },  
  { score: { $meta: "textScore" } })  
.sort({ score: { $meta: "textScore" } })
```

```
{ "_id" : "2", "statement" : "MongoDB is the best", "score" : 1.5 }  
{ "_id" : "4", "statement" : "DynamoDB is the best", "score" : 0.75 }  
{ "_id" : "1", "statement" : "MongoDB is the worst.", "score" : 0.75 }
```




- We have text index built on field on which we want perform full text search
- Also note that
- Success lies in the computation of similarity



More characteristics of Text Indexes

- We can specify language for removing “stop words” and “stemming”.
- A collection can have at most one text index.
- However we can include multiple fields in an index.
- We can also specify weights to each field match
- We can also create “wildcard text” indexes that means index gets built on all text fields,
- So forth



More powerful full text search on Mongo DB

- Full text becomes more important in semi structured data where we find many fields are in text forms.
- [Elasticsearch](#) or [Apache Solr](#) are popular large scale full text search engines with many machine intelligence features.
- MongoDB allows us integrating with such external full text search engines.
- “MongoDB Atlas” already provides integrated solution with lucene

Getting started with MongoDB Atlas Full-Text Search

<https://www.mongodb.com/blog/post/getting-started-with-mongodb-atlas-fulltext-search>



Mongo DB Query Planner

- In database world “Query Optimization” often refers to putting an efficient “**query plan**” for executing a query.
- What we mean by “query plan”?
 - Basically means **having a execution pipeline**.
 - Order of database operations (at logical level): selection (filter), projection, grouping, aggregation, etc.
 - Order and “method” of operations at physical level. Often there are multiple methods of performing a logical operations, and one is more suited over the other in certain situations. For instance “collection scan”, sort-scan, index-scan for searching through a collection.
- Mongo DB also provides an underlying query optimizer.



Mongo DB Query Planner

- For a query, the MongoDB query optimizer chooses and caches the most efficient query plan given the available indexes.
- The evaluation of the most efficient query plan is based on the number of “work units” (works) performed by the query execution plan when the query planner evaluates candidate plans.
- The associated plan cache entry is used for subsequent queries with the same “query shape”.
- To view query plan for a query, we can use `db.collection.explain()` or the `cursor.explain()` . Here we see what plan is winner and its details.



`db.collection.explain()`

- Helps you finding out
 - What plan has been finally executed - what has been various stages?
 - What other query plans
 - Is query using index for sort, and so
 - What part of query



Examples: `db...explain()`

- Here are some examples

```
db.zips.explain().find({ "zip": "35004" })
```

```
db.zips.explain("executionStats").find({ "zip": "35004" } })
```

```
db.zips.explain().find(  
  { "state": "PA", pop: { $gte: 50000 } }, { "_id": 0, "city": 1 }  
)
```

```
db.zips.explain().aggregate( [  
  { $group: { _id: "$state", totalPop: { $sum: "$pop" } } },  
  { $match: { totalPop: { $gte: 10*1000*1000 } } }  
)
```



References and Further Readings

- [1] <https://docs.mongodb.com/manual/>
- [2] Most content put here is from:
<https://docs.mongodb.com/manual/indexes/>
- [3] MongoDB University course “Mongo DB Performance”:
https://university.mongodb.com/mercury/M201/2021_March_16
- [4] Getting started with MongoDB Atlas Full-Text Search
<https://www.mongodb.com/blog/post/getting-started-with-mongodb-atlas-fulltext-search>
- [5] Apache Solr: Because your Database is not a Search: Engine
<https://programmaticponderings.com/2019/02/25/apache-solr-because-your-database-is-not-a-search-engine/>
- [6] Solr vs. Elasticsearch: Who’s The Leading Open Source Search Engine
<https://logz.io/blog/solr-vs-elasticsearch/>