

Introduction to Apache Spark



pm jat @ daiict



What is Spark - recap

- Following are two main revolutionary features that make Spark a amazing solution for cluster computing
 - **“In Memory”, “Distributed”, “Fault Tolerant”** collection of objects (called as RDD)
 - Simple Programming Abstractions;
- Spark offers a revolutionary programming paradigm that makes distributed programming like a desktop programming



Resilient Distributed Dataset (RDD) - recap

- The main abstraction in Spark is **Resilient Distributed Dataset (RDD)**
- RDD are collection of **distributed, fault tolerant “object collection”** partitioned across a set of machines. (Note that “Object collection” here is **“in memory”**)
- A simple analogy; a **“Distributed, Fault Tolerant Array List”**
- RDD objects can explicitly be cached in memory, and reused in consequent calls. This “in memory” processing is what makes Spark, amazing fast!
- For **Fault Tolerance**; RDD objects themselves are not replicated, but maintain information that a partition can be rebuilt if a node fails



Resilient Distributed Dataset (RDD) - recap

- Characteristics of RDD:
 - a read-only,
 - partitioned collection of records.
 - Fault tolerant
 - Lazily evaluated
 - Can be cached
- RDDs can only be created by:
 - (1) by reading from data file,
 - (2) distributing (parallelization) a local collection to multiple nodes,
 - (3) by applying a “transformations” on existing rdd

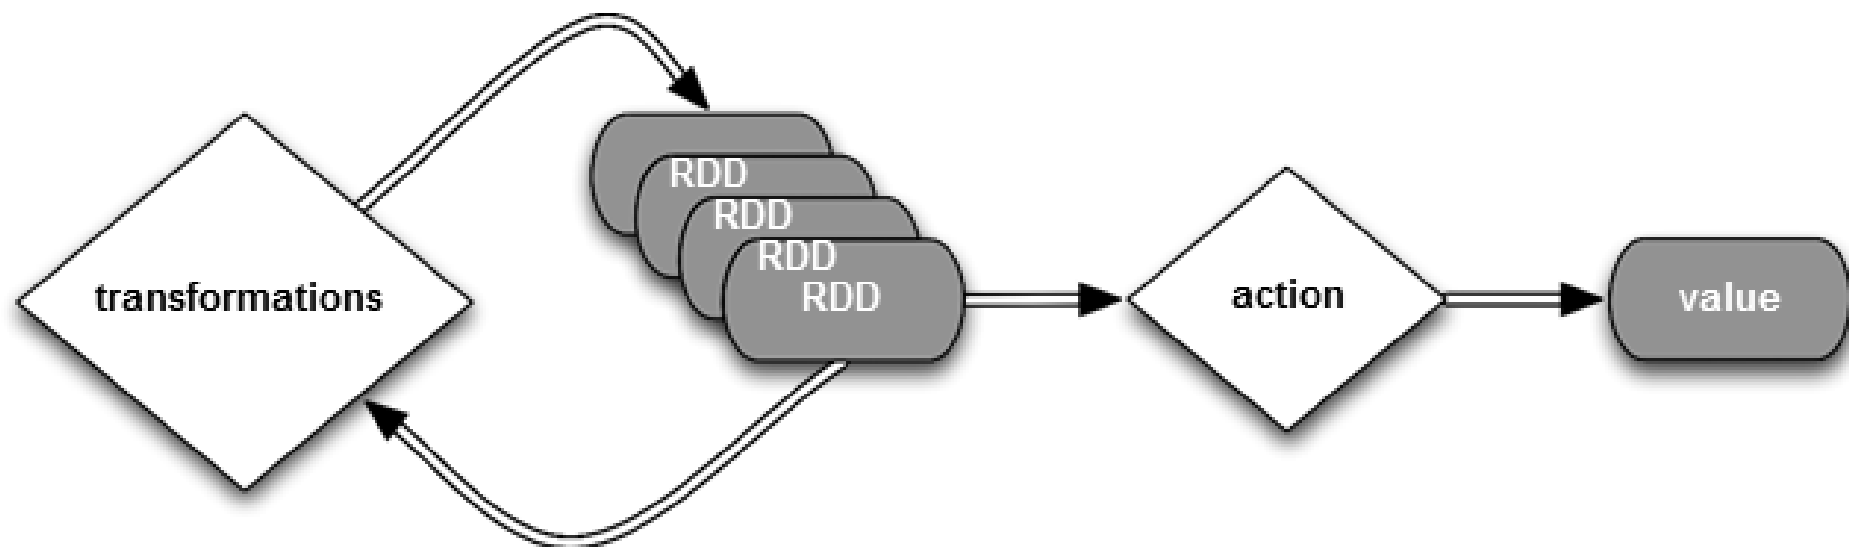


Operations on RDDs

- Operations
 - Transformations, and
 - Actions
- Recall: **RDDs are immutable (read only)**, that is any operation does not change content of an RDD but generates new RDD or some value.



RDD Transformations and Actions





Operations on RDDs

- Transformations
 - map, filter, reduce, group-by, and join
- Actions
 - count (which returns the number of elements in the dataset),
 - collect (which returns the elements themselves), and
 - save (which saves on storage)
- Other operations
 - Persistence/Caching, and
 - “Partitioning”



RDD Transformations^[4]

- Transformations create a new RDD, and are “lazy operations”; i.e. computed only when required!
- Some of the may require **shuffling** of data: like reduceByKey, partition, sort, join etc.

<i>map</i> ($f : T \Rightarrow U$)	:	$\text{RDD}[T] \Rightarrow \text{RDD}[U]$
<i>filter</i> ($f : T \Rightarrow \text{Bool}$)	:	$\text{RDD}[T] \Rightarrow \text{RDD}[T]$
<i>flatMap</i> ($f : T \Rightarrow \text{Seq}[U]$)	:	$\text{RDD}[T] \Rightarrow \text{RDD}[U]$
<i>sample</i> (<i>fraction</i> : Float)	:	$\text{RDD}[T] \Rightarrow \text{RDD}[T]$ (Deterministic sampling)
<i>groupByKey</i> ()	:	$\text{RDD}[(K, V)] \Rightarrow \text{RDD}[(K, \text{Seq}[V])]$
<i>reduceByKey</i> ($f : (V, V) \Rightarrow V$)	:	$\text{RDD}[(K, V)] \Rightarrow \text{RDD}[(K, V)]$
<i>union</i> ()	:	$(\text{RDD}[T], \text{RDD}[T]) \Rightarrow \text{RDD}[T]$
<i>join</i> ()	:	$(\text{RDD}[(K, V)], \text{RDD}[(K, W)]) \Rightarrow \text{RDD}[(K, (V, W))]$
<i>cogroup</i> ()	:	$(\text{RDD}[(K, V)], \text{RDD}[(K, W)]) \Rightarrow \text{RDD}[(K, (\text{Seq}[V], \text{Seq}[W]))]$
<i>crossProduct</i> ()	:	$(\text{RDD}[T], \text{RDD}[U]) \Rightarrow \text{RDD}[(T, U)]$
<i>mapValues</i> ($f : V \Rightarrow W$)	:	$\text{RDD}[(K, V)] \Rightarrow \text{RDD}[(K, W)]$ (Preserves partitioning)
<i>sort</i> ($c : \text{Comparator}[K]$)	:	$\text{RDD}[(K, V)] \Rightarrow \text{RDD}[(K, V)]$
<i>partitionBy</i> ($p : \text{Partitioner}[K]$)	:	$\text{RDD}[(K, V)] \Rightarrow \text{RDD}[(K, V)]$



map transformation

$\text{map}(\text{map-function: } T \rightarrow U): \text{RDD}[T] \Rightarrow \text{RDD}[U]$

- When applied on RDD of type T, it produces RDD of type U by applying *map function*. Mapping is 1:1
- RDD is produced on “same partitions”
- Following map function [Java Code] transforms RDD of String to RDD of String Arrays?

```
JavaRDD<String> lines = sc.textFile("data/employee.csv");  
JavaRDD<String[]> records = lines.map(line -> line.split(", "));
```



```
JavaRDD<String> lines = sc.textFile("data/employee.csv");  
JavaRDD<String[]> records = lines.map(line -> line.split(","));
```

Input: lines [RDD<String>]

```
101, John, 9-Jan-55, M, 30000, 102, 5  
102, Franklin, 8-Dec-45, M, 40000, 105, 5  
103, Joyce, 31-Jul-62, F, 25000, 102, 5  
104, Ramesh, 15-Sep-52, M, 38000, 102, 5  
105, James, 10-Nov-27, M, 55000, , 1  
106, Jennifer, 20-Jun-31, F, 43000, 105, 4  
107, Ahmad, 29-Mar-59, M, 25000, 106, 4  
108, Alicia, 19-Jul-58, F, 25000, 106, 4
```

Output: records [RDD<String[]>]

map

```
<101, John, 9-Jan-55, M, 30000, 102, 5>  
<102, Franklin, 8-Dec-45, M, 40000, 105, 5>  
<103, Joyce, 31-Jul-62, F, 25000, 102, 5>  
<104, Ramesh, 15-Sep-52, M, 38000, 102, 5>  
<105, James, 10-Nov-27, M, 55000, , 1>  
<106, Jennifer, 20-Jun-31, F, 43000, 105, 4>  
<107, Ahmad, 29-Mar-59, M, 25000, 106, 4>  
<108, Alicia, 19-Jul-58, F, 25000, 106, 4>
```



map (producing Pair RDD)

- A variation of RDD, that is, RDD of **<K,V> pairs**, called Pair RDD. Most aggregation operations are performed on such RDDs.
- If we want to produce Pair RDD through Map and often we require doing this, a variation of map function called **mapToPair** is available in Java (some languages like Python, Scala do it implicitly).
- Here is a map example that produces pairs of dno as key and a tuple <salary,1> as value.



map (producing Pair RDD)

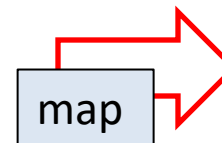
Pseudo Code

```
records = lines.map(line -> line.split(",");  
dnoSalPairs = records.map(  
    record -> (record[6], (int) record = [4])  
);
```

Input: records [RDD<String[]>]

```
<101,John,9-Jan-55,M,30000,102,5>  
<102,Franklin,8-Dec-45,M,40000,105,5>  
<103,Joyce,31-Jul-62,F,25000,102,5>  
<104,Ramesh,15-Sep-52,M,38000,102,5>  
<105,James,10-Nov-27,M,55000,,1>  
<106,Jennifer,20-Jun-31,F,43000,105,4>  
<107,Ahmad,29-Mar-59,M,25000,106,4>  
<108,Alicia,19-Jul-58,F,25000,106,4>
```

Output: **Pair RDD** dnoSalPairs
[String, Tuple<Int,Int>]



key	value
5	<30000,1>
5	<40000,1>
5	<25000,1>
5	<38000,1>
1	<55000,1>
4	<43000,1>
4	<25000,1>
4	<25000,1>



map (producing Pair RDD)

Java Code

```
JavaPairRDD<String, Tuple2<Integer, Integer>> dnoSalPairs  
= records.mapToPair(rec -> new Tuple2<>(  
    rec[6], new Tuple2<>(Integer.parseInt(rec[4]),1))  
);
```

key value

Input: records [RDD<String[]>]

```
<101,John,9-Jan-55,M,30000,102,5>  
<102,Franklin,8-Dec-45,M,40000,105,5>  
<103,Joyce,31-Jul-62,F,25000,102,5>  
<104,Ramesh,15-Sep-52,M,38000,102,5>  
<105,James,10-Nov-27,M,55000,,1>  
<106,Jennifer,20-Jun-31,F,43000,105,4>  
<107,Ahmad,29-Mar-59,M,25000,106,4>  
<108,Alicia,19-Jul-58,F,25000,106,4>
```

Output: **Pair RDD** dnoSalPairs
[String, Tuple2<Int,Int>]

key value

5	<30000,1>
5	<40000,1>
5	<25000,1>
5	<38000,1>
1	<55000,1>
4	<43000,1>
4	<25000,1>
4	<25000,1>

map



map transformation

- Note that instead of two maps for computing $\langle \text{Dno}, \langle \text{Salary}, 1 \rangle \rangle$ pairs, we can do it one map too.
- Shown on next slide
- While second one sounds like shorter, Spark's strategy of lazy evaluation and optimizer should free us from writing either way.



How do following two solutions differ?

- How do following two solutions for computing $\langle \text{Dno}, \langle \text{Salary}, 1 \rangle \rangle$ differ?

```
lines = sc.textFile("data/employee.csv");
records = lines.map(line -> line.split(","));
dnoSalPairs = records.map(
    record -> (record[6], (int) record = [4])
);
```

```
lines = sc.textFile("data/employee.csv");
dnoSalPairs = lines.map( {
    record = line.split(",");
    line -> (record[6], (int) record = [4]);
});
```



filter transformation

`filter(filter-func:T->bool):RDD[T]->RDD[T]`

- Return a new dataset formed by selecting those elements of the source for which *filter-function* returns true [Java Code]

```
JavaRDD<String> lines = sc.textFile("data/log.txt");
JavaRDD<String> debugLines = lines.filter( line -> line.contains("DEBUG") );
System.out.println( "Count - Debug Line: " + debugLines.count() );
List<String> top3 = debugLines.take(3);
top3.forEach( line -> System.out.println(line));
```

- Note the code above also uses two actions “`take`”, and “`count`”



flatMap transformation

`flatMap(mapfunc:T->Seq(U)): RDD[T]->RDD[U]`

- Similar to map, but each input item can be mapped to 0 or more output items,
- so *map-function* that we define should return a Sequence rather than a single item
- Following code transforms a RDD of “lines” to RDD of “words”, though both are happens to be of type String. Type wise both are of String though!

```
JavaRDD<String> words = lines.flatMap(  
    line -> Arrays.asList(line.split(" ")).iterator());
```



reduceByKey transformation

`reduceByKey(reduce-func: (V, V) -> V)`
`: RDD[(K, V)] -> RDD[(K, V)]`

- When called on a dataset of (K, V) pairs (Pair RDD), returns a dataset of (K, V) pairs. where the values for each key are aggregated using the given *reduce-function*.
- The function may also provide additional parameter for number of reduce tasks.
- Here is an example.

```
sums = dnoSalPairs.reduceByKey((x, y) -> x + y);
```



reduceByKey transformation

```
sums = dnoSalPairs.reduceByKey((x, y) -> x + y);
```

Input: Pair RDD

dnoSalPairs[<String, Integer>]

5, 30000
5, 40000
5, 25000
5, 38000
1, 55000
4, 43000
4, 25000
4, 25000

Reduce By Key

output: Pair RDD

sums[<String, Integer>]

4: 93000
5: 133000
1: 55000

```
reduceByKey(reduce-func: (V, V) -> V)  
: RDD[(K, V)] -> RDD[(K, V)]
```



reduceByKey transformation

```
reduceByKey(reduce-func: (V, V) -> V)  
: RDD[(K, V)] -> RDD[(K, V)]
```

Java Code

```
JavaPairRDD<String, Integer> sums  
= dnoSalPairs.reduceByKey((x, y) -> x + y);
```

- Parameters to the function: value-1 (x) and Value-2 (y) each of type Integer. Can read them as following value-1 being the first one (for a distinct key, where as value-2 varies from second to last)
- Computes: performs simply addition, and returns!
- Note about type parameters: K is String, V is Integer



Compute Department wise Total Salary using Reduce-By-Key

- Compute Pipeline: `map(extract dno, sal) → reduceByKey() → result(dno, sumsal)`

```
lines = sc.textFile("data/employee.csv");

dnoSalPairs = lines.mapToPair( {
    record = line.split(",");
    line -> (record[6], (int) record = [4]);
});

sums = dnoSalPairs.reduceByKey((x, y) -> x + y);

//collect and output on console
output = sums.collect();
output.foreach( t -> print(t._1 + ": " + t._2));
```

Pseudo Code



Compute Department wise Total Salary using Reduce-By-Key

Java Code

```
JavaRDD<String> lines = sc.textFile("data/employee.csv");
JavaPairRDD<String, Integer> dnoSalPairs
= lines.mapToPair( line ->
    new Tuple2<>(line.split(",")[6], Integer.parseInt(line.split(",")[4]))
);
JavaPairRDD<String, Integer> sums
= dnoSalPairs.reduceByKey((x, y) -> x + y);
```

MAP: Dno , Sal Pairs

ReduceByKey: Dno , SumSal



Compute Department wise Average Salary using Reduce-By-Key

- Note alone Reduce-By-Key is not enough to compute average; we require additional computation (a map task) after reduce-by-key!
- Here is a pipeline (red-one is additional step w.r.t SUM:
Compute Pipeline:
map(extract dno, <sal,1>)
 → reduceByKey()
 → **map(compute average)**
 → result(dno,avgsal)



Compute Department wise Average Salary using Reduce-By-Key

```
lines = sc.textFile("data/employee.csv");
dnoSalPairs = lines.map( {
    record = line.split(",");
    line -> (record[6], ((int) record[4], record[5]))
});

sums = dnoSalPairs.reduceByKey(
    (v1,v2) -> (v1._1 + v2._1, v1._2 + v2._2)
);

aggregate = sums.map( v -> {
    key = v._1;
    avg = 1.0 * v._2._1/v._2._2;
    return (key, avg);
});
```

MAP:
Dno , Sal Pairs

ReduceByKey:
Dno , SumSal

MAP: Dno , AvgSal



Compute Department wise Average Salary using Reduce-By-Key

- Compute Pipeline: `map(extract dno, <sal,1>)` → `reduceByKey()` → **map(compute average)** → result

```
JavaRDD<String> lines = sc.textFile("data/employee.csv");
JavaPairRDD<String, Tuple2<Integer, Integer>> dnoSalPairs
= lines.mapToPair( line ->
    new Tuple2<>(line.split(",")[6],
        new Tuple2<>(Integer.parseInt(line.split(",")[4]),1))
);

JavaPairRDD<String, Tuple2<Integer, Integer>> sums
= dnoSalPairs.reduceByKey((v1,v2)
    -> new Tuple2<>(v1._1 + v2._1, v1._2 + v2._2));

JavaPairRDD<String, Double> aggregate = sums.mapToPair( (v)
    -> new Tuple2<String, Double>(v._1, 1.0 * v._2()._1/v._2()._2));
```

Java Code

MAP:
Dno , Sal Pairs

ReduceByKey:
Dno , SumSal

MAP: Dno , AvgSal



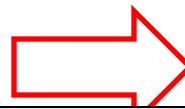
Reduce-By-Key for Average

- For computing Average, function for ReduceByKey is defined such that it produces <salary, sum(count)> pairs for every key (dno).

```
sums = dnoSalPairs.reduceByKey(  
    (v1,v2) -> (v1._1 + v2._1, v1._2 + v2._2)  
);
```

Input: dnoSalPairs[<String,<Int,Int>>]

5, <30000,1>
5, <40000,1>
5, <25000,1>
5, <38000,1>
1, <55000,1>
4, <43000,1>
4, <25000,1>
4, <25000,1>



Reduce By Key

output: sums[String,<Int,Int>>]

4, <93000,3>
5, <133000,4>
1, <55000,1>

```
reduceByKey(reduce-func: (V,V) -> V)  
: RDD[(K,V)] -> RDD[(K,V)]
```



Reduce-By-Key for Average

```
reduceByKey(reduce-func: (V,V) -> V)  
: RDD[(K,V)] -> RDD[(K,V)]
```

```
JavaPairRDD[String, Tuple2<Integer, Integer>> sums  
  = dnoSalPairs.reduceByKey((v1,v2)  
    -> new Tuple2<>(v1._1 + v2._1, v1._2 + v2._2));
```

- Parameters to the function: value-1 and Value-2. each of type <Integer, Integer> tuple.
- Computes: performs pairwise sum, and Returns as tuple!
- Note about type parameters:
K is String, and V is Tuple<Integer, Integer>



Map function computing average!

```
aggregate = sums.map( v -> {  
    key = v._1;  
    avg = 1.0 * v._2._1/v._2._2;  
    return (key, avg);  
});
```

Pseudo Code

input: sums[String,<Int,Int>>]

```
4, <93000,3>  
5, <133000,4>  
1, <55000,1>
```

map

output: aggregate[<String,Double>]

```
4: 31000.0  
5: 33250.0  
1: 55000.0
```

JavaPairRDD[String, Double]

```
aggregate = sums.mapToPair( (v)  
-> new Tuple2<>(v._1, 1.0 * v._2()._1/v._2()._2));
```

Java Code



groupBy transformation

`groupBy(): RDD[(K,V)] -> RDD[(K,seq(V))]`

- When called on a dataset of (K, V) pairs, returns a dataset of (K, Iterable<V>) pairs.
- The function may also provide additional parameter for number of tasks.
- Example:

```
groupedRecords = dnoSalPairs.groupByKey();
```



groupByKey transformation

```
groupedRecords = dnoSalPairs.groupByKey();
```

Input: dnoSalPairs[**RDD<?>**]

Group By Key

output: dnoSalPairs[**RDD<?>**]

```
4: <43000,1>, <25000,1>, <25000,1>
5: <30000,1>, <40000,1>, <25000,1>, <38000,1>
1: <55000,1>
```

```
5, <30000,1>
5, <40000,1>
5, <25000,1>
5, <38000,1>
1, <55000,1>
4, <43000,1>
4, <25000,1>
4, <25000,1>
```



Compute Average using Group-By-Key

```
lines = sc.textFile("data/employee.csv");
dnoSalPairs = lines.map( {
    record = line.split(",");
    line -> (record[6], ((int) record = [4],1));
});

groupedRecords = dnoSalPairs.groupByKey();

aggregate = groupedRecords.map( tuple -> {

    String key = tuple._1;
    values = tuple._2;
    int count = 0;
    double sum = 0;
    for( value : values) {
        sum += value._1;
        count += value._2;
    }
    double avg = sum / count;
    return (key, avg);
});
```



“ImageCounter” implementation

Pseudo Code

```
requests = sc.textFile("data/access_log.txt");

//Define filter function that keeps requests that are for images
imgFilter = request -> {

//Define Map function
imgMap = request -> {

//Apply Filter as per the above function
imgRequests = requests.filter( imgFilter );

//Perform MAP as per the defined in line anonymous function
images = imgRequests.mapToPair( imgMap );

//Apply Reduce Function
counts = images.reduceByKey((x, y) -> x + y);

//collect and output on console
output = sums.collect();
output.forEach( t -> print(t._1 + ": " + t._2));
```


“ImageCounter” implementation

//Define filter function that keeps requests that are for images

```
imgFilter = request -> {  
    String url= request.split(" ")[6];  
    String method= request.split(" ")[5];  
    int pos = url.lastIndexOf(".");  
    if(pos != -1 ) {  
        ext=url.substring(pos+1);  
        if(isImg.contains(ext) && method.equals("\GET"))  
            return TRUE;  
    }  
    return FALSE;  
};
```

Pseudo Code

//Define Map function

```
imgMap = request -> {  
    String url= request.split(" ")[6];  
    int pos = url.lastIndexOf(".");  
    String ext=url.substring(pos+1);  
  
    if(ext.equals("gif")) return ("GIF Images",1);  
    else if(isJPG.contains(ext)) return ("JPG Images",1);  
    else return ("Other Images",1);  
};
```



Writing “functions” to be sent as an argument to Transformations

- Java being strongly typed language makes it bit verbose. Agree Java makes it bit complex!
- Provides following set of interfaces. An object of class implementing one of these interfaces can be sent as an argument to a transformation function.

`Function<T1,R> //for processing RDDs`

Method: `R call(T1 v1)`

`Function2<T1,T2,R> /////for processing Pair RDDs`

Method: `R call(T1 v1, T2 v2)`

`PairFunction<T,K,V> //for returning Pairs`

Method: `Tuple2<K,V> call(T t)`



Examples: “function” parameters to transformations

Function<T1,R>
Method: R call(T1 v1)

Filter Function for ImageFilter

- **imgFilter** gets a new object of anonymous class implementing Function2 interface. The class overrides **Call** method!

```
Function<String, Boolean> imgFilter = new Function<String, Boolean>() {  
    @Override  
    public Boolean call(String request) throws Exception {  
        String url= request.split(" ")[6];  
        String method= request.split(" ")[5];  
        int pos = url.lastIndexOf(".");  
        if(pos != -1 ) {  
            String extension=url.substring(pos+1);  
            if(ImageCounter.isImg.contains(extension) && method.equals("\GET"))  
                return Boolean.TRUE;  
        }  
        return Boolean.FALSE;  
    }  
}
```

Type Parameters to the Function: T1 parameter type, i.e. Parent RDD element type. R is return type



Inline Function for Map task [Java]

PairFunction<T,K,V>

Method: `Tuple2<K,V> call(T t)`

```
//Perform MAP as per the defined in line anonymous function
JavaPairRDD<String, Integer> images = imgRequests.mapToPair(
    new PairFunction<String, String, Integer>() {
        @Override
        public Tuple2<String, Integer> call(String request) throws Exception {
            String url= request.split(" ")[6];
            int pos = url.lastIndexOf(".");
            String extension=url.substring(pos+1);

            if(extension.equals("gif"))
                return new Tuple2<String, Integer>("GIF Images",1);
            else if(ImageCounter.isJPG.contains(extension))
                return new Tuple2<String, Integer>("JPG Images",1);
            else
                return new Tuple2<String, Integer>("Other Images",1);
        }
    });
```

Type Parameters to the PairFunction: T parameter type, i.e. Parent RDD element type. K,V are Key and Value types of Pair to be returned



Implementation of Average using Group-By-Key

```
lines = sc.textFile("data/employee.csv");
dnoSalPairs = lines.map( {
    record = line.split(",");
    line -> (record[6], ((int) record = [4],1))
});

sums = dnoSalPairs.reduceByKey(
    (v1,v2) -> (v1._1 + v2._1, v1._2 + v2._2)
);

aggregate = sums.map( v -> {
    key = v._1;
    avg = 1.0 * v._2._1/v._2._2;
    return (key, avg);
});
```



Implementation of Average using Group-By-Key

Following are RDD Transformations in Average using Group-By-Key

- Read from file ==> RDD **line**[String]
- Map applied (line) ==> RDD **dnoSalPairs**[<String, <Int,Int>>]
- Group By Key applied (dnoSalPairs)
==> RDD **groupedRecords**[<String, Iterable(<Int,Int>)>]
- Map applied to groupedRecords
==> RDD **aggregate** [<String, Double>]

This is desired result!



Java Implementation of Average using Group-By-Key

```
JavaRDD<String> lines = sc.textFile("data/employee.csv");

JavaPairRDD<String, Tuple2<Integer, Integer>> dnoSalPairs
    = lines.mapToPair( line ->
        new Tuple2<>(line.split(",")[6],
            new Tuple2<>(Integer.parseInt(line.split(",")[4]),1))
    );

JavaPairRDD<String, Iterable<Tuple2<Integer, Integer>>>
    groupedRecords = dnoSalPairs.groupByKey();
JavaPairRDD<String, Double> aggregate = groupedRecords.mapToPair(
    { //Implementation of Map function follows }

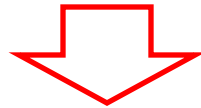
//collect and output on console
List<Tuple2<String, Double>> output = aggregate.collect();
output.forEach( tuple -> System.out.println(tuple._1 + ": " + tuple._2));
```



Map Function for “Computing Aggregate”

Input: `groupedRecords[<String, Iterable[Int]>]`; i.e. DNO, SalPair List

```
4: <43000,1>, <25000,1>, <25000,1>  
5: <30000,1>, <40000,1>, <25000,1>, <38000,1>  
1: <55000,1>
```



output: `aggregate[<String,Double>]`; i.e. DNO, Avg

```
4: 31000.0  
5: 33250.0  
1: 55000.0
```




Map for “Computing Aggregate”

PairFunction<T,K,V>

Method: `Tuple2<K,V> call(T t)`

```
JavaPairRDD<String, Double> aggregate = groupedRecords.mapToPair(  
    new PairFunction<Tuple2<String,  
        Iterable<Tuple2<Integer, Integer>>>, String, Double>() {  
        @Override  
        public Tuple2<String, Double> call(Tuple2<String,  
            Iterable<Tuple2<Integer, Integer>>> tuple)  
            throws Exception {  
            String key = tuple._1;  
            Iterable<Tuple2<Integer, Integer>> values = tuple._2();  
            int count = 0;  
            double sum = 0;  
            for( Tuple2<Integer, Integer> value : values) {  
                sum += value._1;  
                count += value._2;  
            }  
            double avg = sum / count;  
            return new Tuple2<String, Double>(key, avg);  
        }  
    });
```



map function for Computing Average

```
PairFunction<T,K,V>
```

```
Method: Tuple2<K,V> call(T t)
```

- Since we require returning Pair, we use **PairFunction**.
- Parameters to **call** function: Tuple of <DNO, List<Salary, 1>>, i.e. Tuple2<String, Iterable[Tuple2<Integer, Integer>]>
- Computes: Aggregates Value list, and computes $\text{Average} = \Sigma \text{salary} / \Sigma \text{count}$, and
- Return <DNO, Average>
- Type Parameters:
 $T = \text{Tuple2}\langle \text{String}, \text{Iterable}[\text{Tuple2}\langle \text{Integer}, \text{Integer} \rangle] \rangle$,
 $K = \text{String}$, $V = \text{Double}$



aggregateByKey transformation

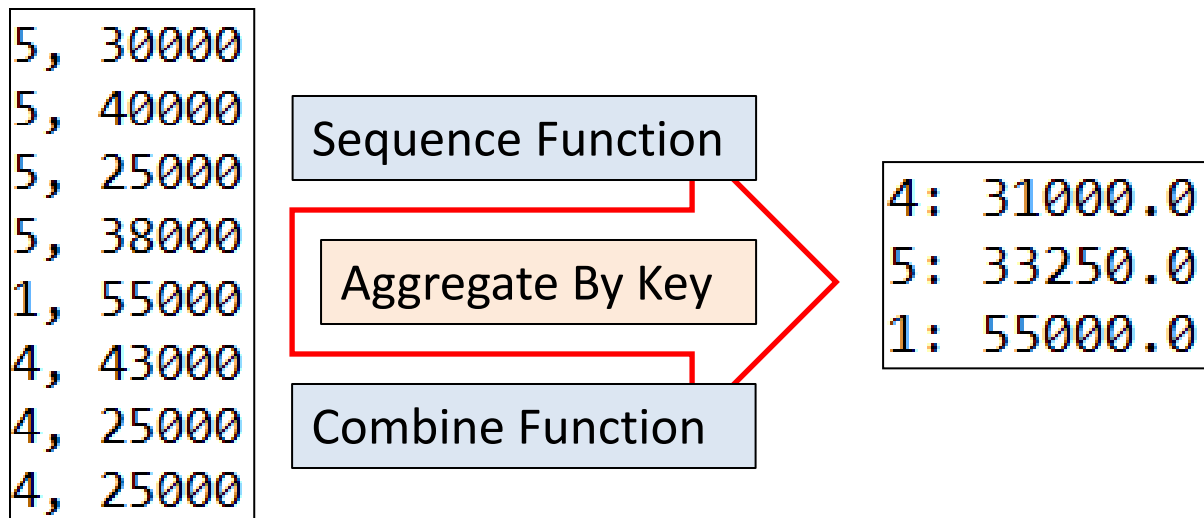
`aggregateByKey(zeroValue, [numPartitions],
seqFunction, combFunction): <K,V> => <K, U>`

- When called on a dataset of (K, V) pairs, returns a dataset of (K, U) pairs where the values for each key are aggregated using the given combine functions and a "zero" value.
- Four Parameters: (1) Zero value (used for initialization purpose while aggregation, (2) optional: number of tasks (3) sequence operation, that runs on every value of value list (sequence), typically accumulates into an accumulator. Corresponds to “**map**” function in a MR job.
(4) combine operations: Corresponds to “**combine**” function in a MR job.
- Allows defining an “aggregated value” type that is different than the input value type. Object of this type is what returned for every key.



Average using Aggregate By Key

- We can use Aggregate By Key for computing





Average using Aggregate By Key

```
lines = sc.textFile("data/employee.csv");
dnoSalPairs = lines.map( {
    record = line.split(",");
    line -> (record[6], ((int) record = [4],1));
});
//define sequence function
addAndCount = ( aggregator, value) -> {
    aggregator.sum += value;
    aggregator.count += 1;
    return aggregator;
};
//define combine function
combine = (aggr1, aggr2) -> {
    aggr1.sum += aggr2.sum;
    aggr1.count += aggr2.count;
    return aggr1;
};
initialAggr = (0,0);
aggregate = dnoSalPairs
    .aggregateByKey(initialAggr, addAndCount, combine);
```



Average using Aggregate By Key

```
JavaRDD<String> lines = sc.textFile("data/employee.csv");

JavaRDD<String[]> records = lines.map(line -> line.split(","));
JavaPairRDD<String, Integer> dnoSalPairs
    = records.mapToPair(rec
        -> new Tuple2<>(rec[6], Integer.parseInt(rec[4])));

Function2<AvgCounter, Integer, AvgCounter> addAndCount = { ... }

Function2<AvgCounter, AvgCounter, AvgCounter> combine = { ... }

AvgCounter initial = new AvgCounter(0,0);
JavaPairRDD<String, AvgCounter> aggregate
    = dnoSalPairs.aggregateByKey(initial, addAndCount, combine);
```



Sequence Function for Average

- Created of type Function2, which is basically a function with two parameters. Try understanding the code below
- Type parameters(U,V,U): U is aggregate type and V is source rdd value type
- Function parameters: 1: aggregate object, 2: value of source rdd
- Function returns: update aggregate object; source value is updated to aggregate object

```
Function2<AvgCounter, Integer, AvgCounter> addAndCount
    = new Function2<AvgCounter, Integer, AvgCounter>() {
    @Override
    public AvgCounter call(AvgCounter a, Integer x) {
        a.sum += x;
        a.count += 1;
        return a;
    }
}
```



Combine Function for Average

- Again of type Function2. Try understanding the code below
- Type parameters(U,U,U): U is aggregate type
- Function parameters: 1: aggregate object a, 2: aggregate object b
- Function returns: updated aggregated object a; b is updated into a.

```
Function2<AvgCounter, AvgCounter, AvgCounter> combine
    = new Function2<AvgCounter, AvgCounter, AvgCounter>() {
    @Override
    public AvgCounter call(AvgCounter a, AvgCounter b) {
        a.sum += b.sum;
        a.count += b.count;
        return a;
    }
};
```




RDD Operations - recap

- Computations in Spark are performed by following two types of operations on RDD
 - Transformations, and
 - Actions
- In all cases RDD is operand. Transformations derive another RDD by applying specified transformation.
- Some transformations are unary while others are binary
- Most Transformations require specifying transformation function
- Actions are used by driver program to collect “some result”!



RDD Operations - summarize

`map(map-function: T->U): RDD[T] => RDD[U]`

- When applied on RDD of type T, it produces RDD of type U by applying *map function*. Mapping is 1:1

`filter(filter-func: T->bool):RDD[T]->RDD[T]`

- Return a new dataset formed by selecting those elements of the source for which *filter-function* returns true [Java Code]

`flatMap(mapfunc:T->Seq(U)): RDD[T]->RDD[U]`

- Similar to map, but each input item can be mapped to 0 or more output items

You can get more concrete coverage at:

<http://spark.apache.org/docs/latest/rdd-programming-guide.html>

<https://training.databricks.com/visualapi.pdf>



RDD Operations - summarize

`reduceByKey(reduce-func: (V,V) -> V)`
`: RDD[(K,V)] -> RDD[(K,V)]`

- When called on a dataset of (K, V) pairs (Pair RDD), returns a dataset of (K, V) pairs. where the values for each key are aggregated using the given *reduce-function*.

`groupByKey(): RDD[(K,V)] -> RDD[(K, seq(V))]`

- When called on a dataset of (K, V) pairs, returns a dataset of (K, Iterable<V>) pairs.

`aggregateByKey(zeroValue, seqFunction,`
`combFunction): <K,V> => <K, U>`

- When called on a dataset of (K, V) pairs, returns a dataset of (K, U) pairs where the values for each key are aggregated using the given combine functions.



RDD Operations - summarize

union(otherRDD): (RDD[T], RDD[T]) => RDD[T]

- Return a new dataset that contains the union of the elements in the source dataset and the argument.

join(otherRDD)

: (RDD[(K, V)], RDD[(K, W)]) => RDD[(K, (V, W))]

- When called on datasets of type (K, V) and (K, W), returns a dataset of (K, (V, W)) pairs with all pairs of elements for each key.

mapValues(map-func: V => W): RDD[K, V] => RDD[K, W]

- Key remains unchanged, only value is transformed. It becomes efficient, partition does not change.



RDD Operations - summarize

`COALESCE(numPartitions, shuffle_flag)`

- Return a new RDD which is reduced to a smaller number of partitions
- Shuffle_flag tells, if shuffling is to be done, or merely merging of some of partitions!

You can get more concrete coverage at:

<http://spark.apache.org/docs/latest/rdd-programming-guide.html>

<https://training.databricks.com/visualapi.pdf>



RDD Transformations^[4]

- Transformations create a new RDD, and are “lazy operations”; i.e. computed only when required!

<i>map</i> (<i>f</i> : <i>T</i> ⇒ <i>U</i>)	:	<i>RDD</i> [<i>T</i>] ⇒ <i>RDD</i> [<i>U</i>]
<i>filter</i> (<i>f</i> : <i>T</i> ⇒ <i>Bool</i>)	:	<i>RDD</i> [<i>T</i>] ⇒ <i>RDD</i> [<i>T</i>]
<i>flatMap</i> (<i>f</i> : <i>T</i> ⇒ <i>Seq</i> [<i>U</i>])	:	<i>RDD</i> [<i>T</i>] ⇒ <i>RDD</i> [<i>U</i>]
<i>sample</i> (<i>fraction</i> : <i>Float</i>)	:	<i>RDD</i> [<i>T</i>] ⇒ <i>RDD</i> [<i>T</i>] (Deterministic sampling)
<i>groupByKey</i> ()	:	<i>RDD</i> [(<i>K</i> , <i>V</i>)] ⇒ <i>RDD</i> [(<i>K</i> , <i>Seq</i> [<i>V</i>])]
<i>reduceByKey</i> (<i>f</i> : (<i>V</i> , <i>V</i>) ⇒ <i>V</i>)	:	<i>RDD</i> [(<i>K</i> , <i>V</i>)] ⇒ <i>RDD</i> [(<i>K</i> , <i>V</i>)]
<i>union</i> ()	:	(<i>RDD</i> [<i>T</i>], <i>RDD</i> [<i>T</i>]) ⇒ <i>RDD</i> [<i>T</i>]
<i>join</i> ()	:	(<i>RDD</i> [(<i>K</i> , <i>V</i>)], <i>RDD</i> [(<i>K</i> , <i>W</i>)]) ⇒ <i>RDD</i> [(<i>K</i> , (<i>V</i> , <i>W</i>))]
<i>cogroup</i> ()	:	(<i>RDD</i> [(<i>K</i> , <i>V</i>)], <i>RDD</i> [(<i>K</i> , <i>W</i>)]) ⇒ <i>RDD</i> [(<i>K</i> , (<i>Seq</i> [<i>V</i>], <i>Seq</i> [<i>W</i>]))]
<i>crossProduct</i> ()	:	(<i>RDD</i> [<i>T</i>], <i>RDD</i> [<i>U</i>]) ⇒ <i>RDD</i> [(<i>T</i> , <i>U</i>)]
<i>mapValues</i> (<i>f</i> : <i>V</i> ⇒ <i>W</i>)	:	<i>RDD</i> [(<i>K</i> , <i>V</i>)] ⇒ <i>RDD</i> [(<i>K</i> , <i>W</i>)] (Preserves partitioning)
<i>sort</i> (<i>c</i> : <i>Comparator</i> [<i>K</i>])	:	<i>RDD</i> [(<i>K</i> , <i>V</i>)] ⇒ <i>RDD</i> [(<i>K</i> , <i>V</i>)]
<i>partitionBy</i> (<i>p</i> : <i>Partitioner</i> [<i>K</i>])	:	<i>RDD</i> [(<i>K</i> , <i>V</i>)] ⇒ <i>RDD</i> [(<i>K</i> , <i>V</i>)]

Cogroup is also called
groupWith

Partition is basically Repartition

You can get more concrete coverage at:

<http://spark.apache.org/docs/latest/rdd-programming-guide.html>

<https://training.databricks.com/visualapi.pdf>



RDD Actions^[4]

- Actions launch a computation to return a value to the driver program or write data to external storage.

<i>count()</i>	:	$\text{RDD}[T] \Rightarrow \text{Long}$
<i>collect()</i>	:	$\text{RDD}[T] \Rightarrow \text{Seq}[T]$
<i>reduce</i> ($f : (T, T) \Rightarrow T$)	:	$\text{RDD}[T] \Rightarrow T$
<i>lookup</i> ($k : K$)	:	$\text{RDD}[(K, V)] \Rightarrow \text{Seq}[V]$ (On hash/range partitioned RDDs)
<i>save</i> ($path : \text{String}$)	:	Outputs RDD to a storage system, <i>e.g.</i> , HDFS

- Reduce action returns a single value of type T. The reduce function should be commutative and associative so that it can be computed correctly in parallel.



RDD Actions^[5]

- More Actions:
 - `take(n)` //if collect can return huge list
 - `takeSample(num, seed)`
 - `takeOrdered(n, ordering_function)`
 - `countByKey`
 - `foreach`



Action Examples

- Code below shows “take”, “count” actions. Should be self explanatory.

```
JavaRDD<String> lines = sc.textFile("data/log.txt");
JavaRDD<String> debugLines = lines.filter( line -> line.contains("DEBUG") );

System.out.println( "Count - Debug Line: " + debugLines.count() );

List<String> top3 = debugLines.take(3);

top3.forEach( line -> System.out.println(line));
```



Persistence and caching of RDD

- We can make RDDs to live only in primary memory, or on disk
- RDD API provides two methods for this `persist`, and `cache`.
- `Cache` tells RDD to be living only in primary memory, where as
- With `persist`, we can specify various methods of persistence.
Typical values: `MEMORY_ONLY`, `MEMORY_AND_DISK`, `DISK_ONLY`, `MEMORY_ONLY_SER`, and so.



RDD – lazy transformations

- RDD transformations are “**lazy operations**”, i.e. they are evaluated on request of some “Action”
- RDD execution engine maintains “**lineage graph**”; all operations specified on RDD are added to lineage graph.
- Lineage graph is evaluated on demand.
- This arrangement helps in optimizing the execution
 - Multiple operations can be performed in a single scan
 - Amount of Data to be shuffled can be minimized by performing local aggregations, etc



Example: Lineage Graph^[4]

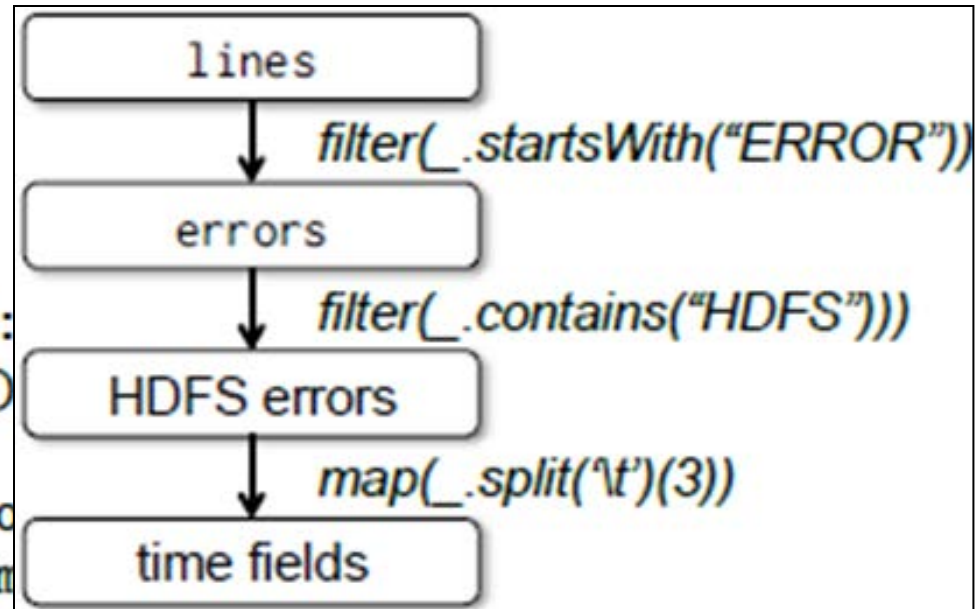
```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
errors.persist()
```

```
errors.count()
```

```
// Count errors mentioning MySQL:
errors.filter(_.contains("MySQL"))
        .count()
```

```
// Return the time fields of error
// HDFS as an array (assuming time
// number 3 in a tab-separated format):
```

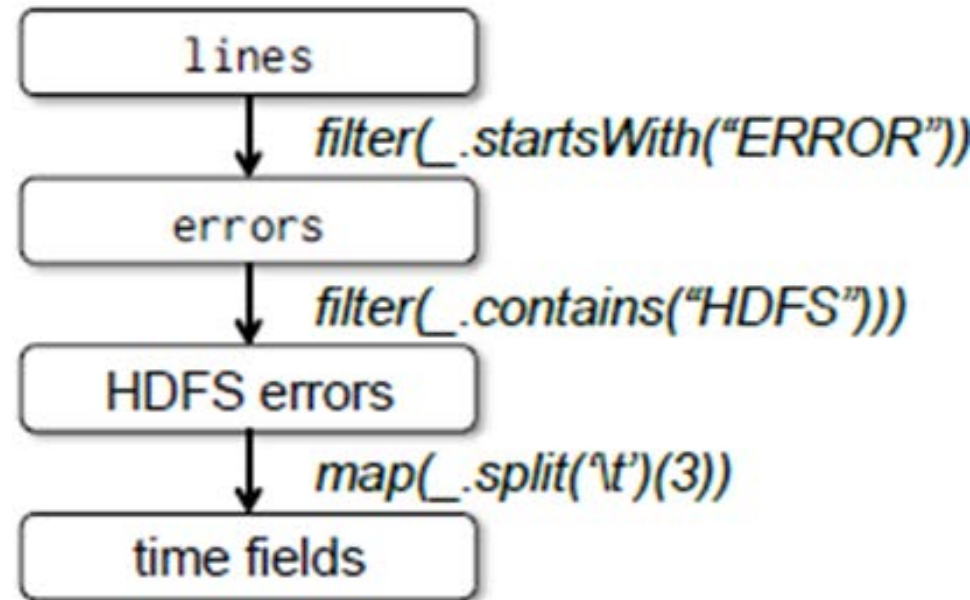
```
errors.filter(_.contains("HDFS"))
        .map(_.split('\t')(3))
        .collect()
```





RDDs – fault tolerance^[4]

- Replication is mechanism of fault tolerance.
- RDD themselves are not replicated but they are reconstructed on another node if a computing node holding part of an RDD fails.
- Suppose one of the node containing **errors** RDD fails.
- That partition of errors can be computed from other replica of data chunk on other available node





Shared variables

- Shared variables is a mechanism of capturing notion of “Global Variables” – global across computing nodes
- Are of two types
 - **Broadcast variables:** arrangement of making some data always data available on (data) partitions. Broadcast variables are in Read Only Mode.
 - **Accumulators:** a global short of data that tasks can accumulate some counters or aggregators.



Broadcast variables:

- Broadcast variables allow the programmer to keep a read-only variable cached on each machine rather than shipping a copy of it with tasks.
- If a large read-only piece of data (e.g., a lookup table) is used in multiple parallel operations, it is preferable to distribute it to the workers, only once.
- Spark uses sophisticated broadcast algorithms to reduce the communication cost!

Example in Java

Broadcast

```
Broadcast<int[]> broadcastVar = sc.broadcast(new int[] {1, 2, 3});
```

Accessing in some local function like map or so

```
broadcastVar.value();
```

```
// returns [1, 2, 3]
```

<http://spark.apache.org/docs/latest/rdd-programming-guide.html>



Accumulators

- Accumulators are variables that are only “added” through an associative and commutative operation and can therefore be efficiently supported in parallel.
- Accumulators typically allows, mappers to put data in parallel!
- Therefore, a mechanism of building some aggregations, like sums and counters (as in Map Reduce) .
- Spark natively supports accumulators of numeric types, and programmers can add support for new types.



Accumulators

- Java example

```
LongAccumulator accum = jsc.sc().longAccumulator();
```

Updaing in some local function like map, or foreach

```
sc.parallelize(Arrays.asList(1, 2, 3, 4)).foreach(x -> accum.add(x));
```

```
// ...
```

```
// 10/09/29 18:41:08 INFO SparkContext: Tasks finished in 0.317106 s
```

Collecting in Driver code

```
accum.value();
```

```
// returns 10
```



Accumulators

- Typically, it works as following (local accumulation and then aggregation on request)

Accumulators

Accumulable	Value
counter	45

Tasks

Index ▲	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Accumulators	E
0	0	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2016/04/21 10:10:41	17 ms			
1	1	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2016/04/21 10:10:41	17 ms		counter: 1	
2	2	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2016/04/21 10:10:41	17 ms		counter: 2	
3	3	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2016/04/21 10:10:41	17 ms		counter: 7	
4	4	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2016/04/21 10:10:41	17 ms		counter: 5	
5	5	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2016/04/21 10:10:41	17 ms		counter: 6	
6	6	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2016/04/21 10:10:41	17 ms		counter: 7	
7	7	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2016/04/21 10:10:41	17 ms		counter: 17	



Further Reading

- Chapter 3 and Chapter 4 of book “Learning spark: lightning-fast big data analysis”, O'Reilly Media, Inc.", 2015.
 - The book discusses Spark programming in three languages: Scala, Python, Java!
- RDD Programming Guide
<http://spark.apache.org/docs/latest/rdd-programming-guide.html>



Further Higher Abstractions

- RDD are still low level to perform analytical tasks!
- Spark provides higher abstractions
 - Dataframe API
 - Spark-SQL
- These abstractions makes “cluster programming” amazingly simple, and
- Primarily the reason Spark is becoming popular for big data processing.



References

- [1] Zaharia, Matei, et al. "Spark: Cluster computing with working sets." *HotCloud* 10.10-10 (2010): 95.
- [2] Zaharia, Matei, et al. "Apache spark: a unified engine for big data processing." *Communications of the ACM* 59.11 (2016): 56-65.
- [3] Shi, Juwei, et al. "Clash of the titans: Mapreduce vs. spark for large scale data analytics." *Proceedings of the VLDB Endowment* 8.13 (2015): 2110-2121.
- [4] Zaharia, Matei, et al. "Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing." *Proceedings of the 9th USENIX conference on Networked Systems Design and Implementation*. USENIX Association, 2012.
- [5] RDD Programming Guide:
<http://spark.apache.org/docs/latest/rdd-programming-guide.html>