

Spark Dataframe API and Spark SQL

Spark  *SQL*



pm jat @ daiict



Further Higher Abstractions

- RDD are still “low level” to perform analytical tasks!
- Spark provides higher level abstractions
 - “Dataframe” API (more “**structured RDDs**”)
 - Spark-SQL (SQL interface)
- These abstractions makes “cluster programming” amazingly simple, and
- Primarily the reason Spark is becoming popular for big data processing.



SQL Interface over huge raw files!

- We require SQL interface over raw data file, so that, we are able to run queries without loading data into some database systems
 - For many ETL operations, load time into database systems is forbiddingly high
 - Schema check happens only when we run the query (Read time schema validation)
- This makes running ad-hoc queries on data files quick
- This has been motivation for Pig, Hive, and now Spark-SQL



Before Spark SQL

- Hive from Facebook
- Pig from Yahoo
 - Not really SQL but a scripting language like perl
- Cloudera Impala
- Google Dremel



Hive and Pig

- **Hive:** SQL like interface for HDFS files. Initially developed at Facebook (now Apache project).

facebook



```
SELECT count(*) FROM users
```

In reality, 90+% of MR jobs are generated by Hive SQL

- **Pig:** scripting language for various data transformations. Initially developed at Yahoo. Now again apache project

YAHOO!



```
A = load 'foo';  
B = group A all;  
C = foreach B generate COUNT(A);
```

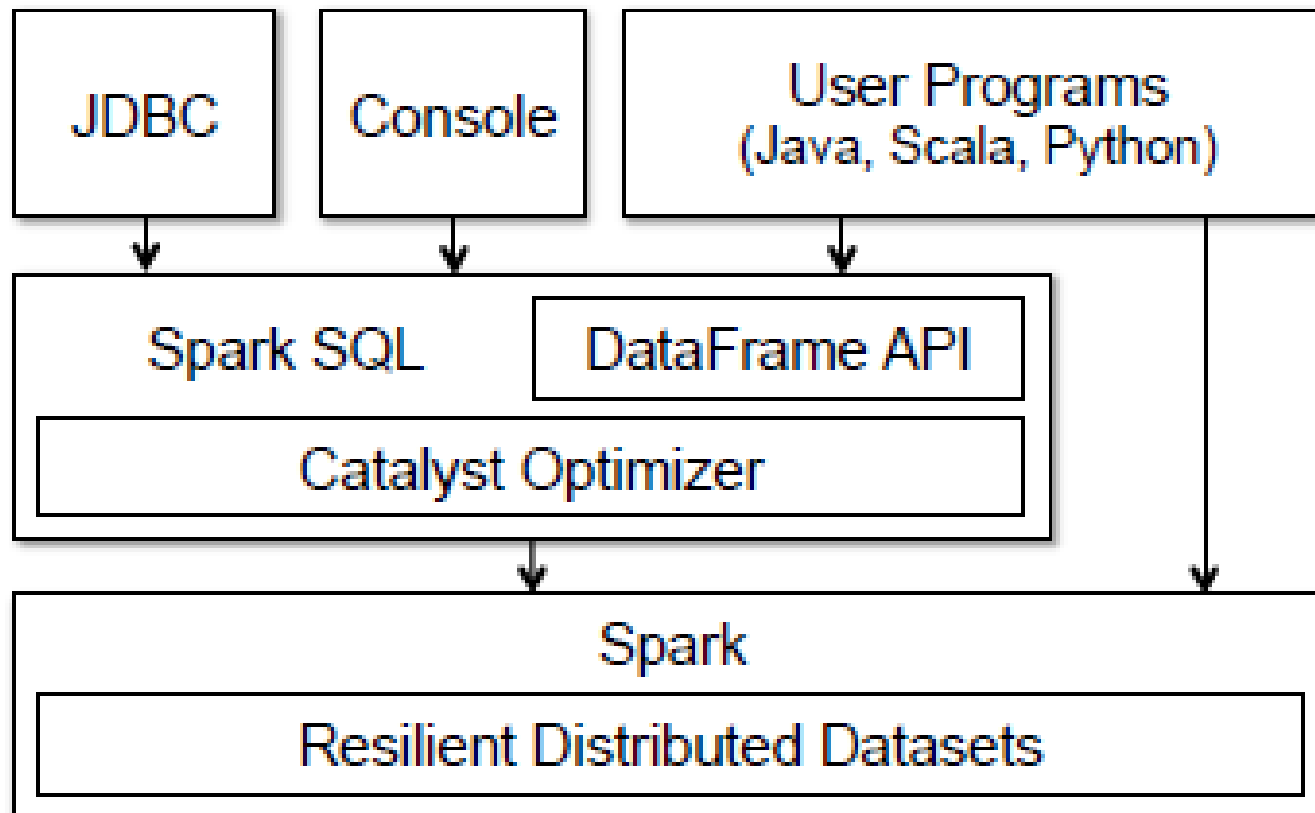


Spark-SQL

- Spark SQL was introduced in 2015 through paper “[Spark SQL: Relational data processing in spark.](#)” in ACM SIGMOD
- Initial effort to have SQL interface over Spark was **Shark**, and was basically an adaption of “Hive over Spark”.
- Shark, however
 - could not integrate well will with intermediate RDD datasets, and
 - Secondly Hive optimizer could not properly adapted to Spark as it was primarily designed for Map Reduce.
- And, we have [Spark-SQL](#)[1, 2015]



Spark-SQL – Stack^[1]



[1] Armbrust, Michael, et al. "Spark sql: Relational data processing in spark." *Proceedings of the 2015 ACM SIGMOD international conference on management of data*. ACM, 2015



Spark-SQL - key characteristics

- Spark-SQL has following key characteristics-
 - Performance (since underlying engine is Spark)
 - Ability to access “procedures” from SQL statements
 - Opens up gate (for SQL) to interact with ML, Graph Processing algorithms as “User Defined Functions”
 - Query Optimizer
 - `Select regmodel.predictsalary(resume) from employee;`



Spark SQL - Programming Interface

- Spark SQL runs as a library on top of Spark (refer stack diagram)
- It has become buzzword for “**Declarative Big Data Processing**”
 - Write less code
 - Optimize the execution
- Spark exposes SQL interfaces, which can be accessed through
 - Command Line Interface (Console)
 - DataFrame API integrated into Spark programs
 - JDBC/ODBC access (through Spark Thrift Server)



Spark “Data Frame API” Source: [1]

- The main abstraction in Spark SQL’s API is a **DataFrame**, a “distributed collection of rows with a **homogeneous schema**”.
- A DataFrame is equivalent to a table in a relational database, and can also be manipulated in similar ways.
- Unlike RDDs, DataFrames keep track of their schema and support various relational operations on them.
- DataFrames can be constructed from data files, tables in a system catalog (based on external data sources) or from existing RDD objects!



Transformations and Actions on DataFrame

Source: [2]

Transformation examples

- filter
- select
- drop
- intersect
- join

Action examples

- count
- collect
- show
- head
- take

- *Note: Dataframe evaluations are lazy. Transformations are just getting expressed (not executed immediately). It is the request of actions that lead to real execution (and this is supported by optimization)*



DataFrame transformation are Lazy

Source: [1]

- Unlike traditional (like in Panda and R) data frame APIs, Spark Data Frames are lazy,
 - in spark, each DataFrame object represents a logical plan to compute a dataset frame
- It is like, when specified, it adds to execution path, and executed only when some action is to be performed.
- This enables optimization across all operations that were used to build the Data Frame.



Data Frames and RDDs

- DataFrames are built on top of the RDD and are “structured”, that is have schema attached with them.
- We can call them Schema aware RDDs - element type is “Row with Schema”
- Also immutable
- DataFrames are more structured, provides abstraction similar to a database table. This makes it very simple to program.
- Dataframes are promised to be more efficient, because operations on them **can be optimized through “Catalyst”, the optimizer for Spark-SQL!**



Programming Dataframes



Spark Session

- Prior to Spark 2.0, Spark had three Contexts
 - Spark Context: for RDD operations
 - SQL Context: for Spark SQL operations
 - Hive Context: for Hive queries
- However Spark 2.0 onwards, there is a unified context called Spark Session, all operations can be performed through spark session, therefore
- Spark Session is more recent!



Initialize Spark Session

- Before doing anything on Spark Dataframe API or Spark SQL we require creating Spark Session
- Typically done as following

```
SparkSession spark = SparkSession.builder()  
    .appName("Spark Demo")  
    .master("local")  
    .getOrCreate();
```




Reading from a CSV file as Dataframe

- Typically done as following:

```
Dataset<Row> emp = spark.read()  
    .option("header", "true")  
    .option("sep", ",")  
    .csv("data/employee_m.csv");
```

- Peek into read data rows Can from the CSV file.

```
//show all rows
```

```
emp.show(); //select * from emp;
```

```
//show first 3 rows
```

```
emp.show(3); //select * from emp limit 3;
```



Reading from a CSV file as Dataframe

- By default
 - If No Header Row in data, then columns are named as _c0, _c1, _c2, and so on
 - To indicate that file has header row, we say:
`.option("header", "true")`
 - Schema remains undefined. By default data type for each column remains string.
- We can ask spark to automatically infer the schema, or we can manually supply the schema before reading a CSV file.



Inferring Schema

- Set : `.option("inferSchema", "true");` to auto infer the schema (in read options)
- We can dump the inferred schema as following:
`emp.printSchema();`

```
root
|-- eno: integer (nullable = true)
|-- name: string (nullable = true)
|-- dob: timestamp (nullable = true)
|-- gender: string (nullable = true)
|-- salary: integer (nullable = true)
|-- sup_eno: integer (nullable = true)
|-- dno: integer (nullable = true)
```

- For inferring the schema, processor require scanning full file. May be expensive when file is huge, and may not always be successful.



Manually Specifying the Schema

(1) Define Schema:

```
StructType schema = DataTypes.createStructType(new StructField[] {  
    //eno,name,dob,gender,salary,sup_eno,dno  
    DataTypes.createStructField("eno", DataTypes.StringType, false),  
    DataTypes.createStructField("name", DataTypes.StringType, true),  
    DataTypes.createStructField("dob", DataTypes.DateType, true),  
    DataTypes.createStructField("gender", DataTypes.StringType, true),  
    DataTypes.createStructField("salary", DataTypes.IntegerType, true),  
    DataTypes.createStructField("sup_eno", DataTypes.StringType, true),  
    DataTypes.createStructField("dno", DataTypes.IntegerType, true)  
});  
//System.out.println(schema.prettyJson());
```

(2) Specify Schema:

```
Dataset<Row> emp = spark.read()  
    .option("header", "true")  
    .option("sep", ",")  
    .schema(schema)  
    .csv("data/employee_m.csv");
```



Reading from a JSON file as Dataframe

- Data frame object is constructed and used as following:

```
Dataset<Row> emp1 = spark.read()  
    .json("data/employees.json");  
emp1.printSchema();  
emp1.show();
```

name	salary
Michael	3000
Andy	4500
Justin	3500
Berta	4000

```
root  
 |-- name: string (nullable = true)  
 |-- salary: long (nullable = true)
```



Examples – Filter and Sort

Filter (/where)

```
emp.filter( "salary > 30000" ).show();
```

```
emp.where( "salary > 40000" ).show();
```

```
emp.filter( emp.col("salary").geq(40000)  
.and(emp.col("dno").equalTo(5))).show(); //More TYPE SAFE
```

Sort/Order By

```
emp.sort(emp.col("dno"),emp.col("salary").desc())  
.show();
```

```
emp.orderBy(emp.col("dno"),emp.col("salary"))  
.show();
```



Examples - Project

Project

```
emp.select("eno", "name", "salary")  
  .where("dno == 4")  
  .show(); //SELECT eno, name, salary from emp  
           //where dno=4
```

```
emp.select(  
  emp.col("eno"),  
  emp.col("name"),  
  emp.col("salary").multiply(1.1)  
) .where("dno == 4")  
  .show();
```



Examples - Aggregation

Aggregation and Group By

```
//on whole Table
```

```
emp.agg(avg(emp.col("salary")),max(emp.col("salary")))  
.show();
```

```
//Group By and Aggregation
```

```
emp.groupBy(emp.col("dno"))  
.agg(avg(emp.col("salary")), max(emp.col("salary")))  
.show();
```




Examples - Join

//Join and Project

```
Dataset<Row> empdep =
```

```
emp.join(dep, emp.col("dno").equalTo(dep.col("dno")));
```

```
empdep.select(emp.col("name"), dep.col("name"),  
              emp.col("salary"))  
.show();
```

//Join and Aggregate

```
emp.join(dep, emp.col("dno").equalTo(dep.col("dno")))  
  .groupBy(dep.col("name"))  
  .agg(sum(emp.col("salary")))  
.show();
```



[Python] Constructing Data Frame

Source: [2]

```
# Construct a DataFrame from a "users" table in Hive.  
df = sqlContext.table("users")  
  
# Construct a DataFrame from a log file in S3.  
df = sqlContext.load("s3n://someBucket/path/to/data.json", "json")
```





[Python] Operation on Data Frames

Source: [2]

```
# Create a new DataFrame that contains only "young" users
young = users.filter(users["age"] < 21)

# Alternatively, using a Pandas-like syntax
young = users[users.age < 21]

# Increment everybody's age by 1
young.select(young["name"], young["age"] + 1)

# Count the number of young users by gender
young.groupBy("gender").count()

# Join young users with another DataFrame, logs
young.join(log, logs["userId"] == users["userId"], "left_outer")
```





[Python] Example: DataFrame API

- Let us say a tab separated data file called “**SalesProduct.txt**”, where attributes Name, and Weight at 2nd and 8th position respectively.
- Following Spark program (using dataframe API) lists Name and Weight of top 15 products in the descending order weight!

```
sqlContext = SQLContext(sc)
```

```
rdd1 = (content.filter(lambda line: line.split("\t")[7] != "NULL")  
        .map(lambda line: (line.split("\t")[1], float(line.split("\t")[7])))  
        )
```

```
df = sqlContext.createDataFrame(rdd1, schema = ["Name", "Weight"])
```

```
df.orderBy("weight", ascending = False).show(15, truncate = False)
```



Spark SQL



Spark SQL

- Spark SQL allows you to manipulate distributed data with SQL queries.
- That is we can execute an SQL statement on a Dataframe object. Though this requires us registering a dataframe as View using `createOrReplaceTempView()`
- Dataframe object provides SQL method for executing SQL statements. This method always returns a dataframe object.
- We can mix DataFrame methods and SQL queries in the same code.



Spark SQL

- Concepts are very similar to relational “Table” and “SQL”

Tables

- In programmer’s perspective table is identical to relational table.
- Here tables can be temporary or be stored.
- We typically have one data file for one table.
- Table data may be stored in “storage formats” that helps in efficient execution of queries.



Spark SQL

- Spark SQL provides same set of commands **DDL + DML**
- However commands do have some additional parameters that are specific to distributed data and clustered computing
 - For instance partitioning, shuffling, and sorting strategy, etc.
- Spark SQL run-time system provides “Spark SQL Engine” that “efficiently” executes Spark-SQL operations.
- SQL statements can be executed from CLI (command level interface) or host programs.
- CLI allows executing SQL commands interactively.



Spark SQL statements

- CREATE TABLE - has options like PARTITION BY, CLUSTER BY
 - Also ALTER TABLE, DROP TABLE, etc
- INSERT, UPDATE, LOAD data,
- CREATE, DROP FUNCTION
- Various commands to add, merge, delete in “data lakes”. Data lake is basically large scale data repository. One can be spark based data lake with SQL interface.
- Here is SQL reference manual from Databrick
<https://docs.databricks.com/spark/latest/spark-sql/index.html>



CREATE table

```
CREATE [TEMPORARY] TABLE [IF NOT EXISTS] [db_name.]table_name
  [(col_name1 col_type1 [COMMENT col_comment1], ...)]
  USING datasource
  [OPTIONS (key1=val1, key2=val2, ...)]
  [PARTITIONED BY (col_name1, col_name2, ...)]
  [CLUSTERED BY (col_name3, col_name4, ...) INTO num_buckets BUCKETS]
  [LOCATION path]
  [COMMENT table_comment]
  [TBLPROPERTIES (key1=val1, key2=val2, ...)]
  [AS select_statement]
```



CREATE table options

- USING
 - Data file storage format
 - Can be one of - TEXT, CSV, JSON, JDBC, PARQUET, ORC, HIVE, DELTA, and LIBSVM
- Option PARTITIONED BY
 - Partition the created table by the specified columns.
- Option CLUSTERED BY
 - Each partition in the created table will be split into a fixed number of buckets by the specified columns. This is typically used within partition for minimizing read time.



SELECT statement

```
SELECT [hints, ...] [ALL|DISTINCT] named_expression[, named_expression, ...]  
  FROM relation[, relation, ...]  
  [lateral_view[, lateral_view, ...]]  
  [WHERE boolean_expression]  
  [aggregation [HAVING boolean_expression]]  
  [ORDER BY sort_expressions]  
  [CLUSTER BY expressions]  
  [DISTRIBUTE BY expressions]  
  [SORT BY sort_expressions]  
  [WINDOW named_window[, WINDOW named_window, ...]]  
  [LIMIT num_rows]
```



SELECT statement

- Option DISTRIBUTE BY
 - Repartition rows in the relation based on a set of expressions. Rows with the same expression values will be hashed to the same worker
- Option CLUSTER BY (WITHIN Partition)
 - Repartition rows in the relation based on a set of expressions and sort the rows in ascending order based on the expressions.
- Also have commands like SAMPLE, ROLLUP, CUBE, etc



Spark SQL Examples

- Suppose, we already have a dataset object `emp` constructed and read some data.
- Below is how we can execute a SQL statement on this. Details of result dataset are printed, and output is shown here.

```
emp.createOrReplaceTempView("employee");  
String sql = "SELECT eno, name, salary FROM employee WHERE dno=4";  
Dataset<Row> emp_dno4 = spark.sql( sql );  
emp_dno4.show();  
emp_dno4.printSchema();
```

```
+---+-----+-----+  
|eno|    name|salary|  
+---+-----+-----+  
|106|Jennifer| 43000|  
|107|   Ahmad| 25000|  
|108|  Alicia| 25000|  
+---+-----+-----+
```

```
root  
|-- eno: string (nullable = true)  
|-- name: string (nullable = true)  
|-- salary: integer (nullable = true)
```



Example: Registering Tables

```
SparkSession spark = SparkSession.builder()  
    .appName("Spark Demo").master("local").getOrCreate();  
  
Dataset<Row> emp = spark.read()  
    .option("header", "true")  
    .option("sep", ",")  
    .option("inferSchema", "true")  
    .csv("data/employee_m.csv");  
  
Dataset<Row> dep = spark.read()  
    .option("header", "true")  
    .option("sep", ",")  
    .option("inferSchema", "true")  
    .csv("data/department_m.csv");  
  
emp.createOrReplaceTempView("employee");  
dep.createOrReplaceTempView("department");
```



Spark SQL Examples

- Various SQL operations on Employee and Department tables.
Code should be self explanatory!

```
String sql = "SELECT d.name, e.name FROM employee e JOIN "  
            + "department d ON (e.dno=d.dno)";  
spark.sql( sql ).show();
```

```
sql = "SELECT e.name, s.name FROM employee e LEFT JOIN "  
      + "employee s ON (e.sup_eno=s.eno)";  
spark.sql( sql ).show();
```

```
sql = "SELECT dno, sum(salary) as TotalSal FROM employee "  
      + "group by dno order by sum(salary) desc ";  
spark.sql( sql ).show();
```




(Python) Example: Spark SQL

```
sqlContext = SQLContext(sc)
```

```
rdd1 = (content.filter(lambda line: line.split("\t")[7] != "NULL")  
        .map(lambda line: (line.split("\t")[1], float(line.split("\t")[7])))  
        )
```

```
df = sqlContext.createDataFrame(rdd1, schema = ["Name", "Weight"])
```

```
df.createOrReplaceTempView("df_table")
```

```
sqlContext.sql(" SELECT * FROM df_table  ORDER BY Weight DESC limit 15").show()
```



Spark SQL – Procedural integration^[1]

- Spark SQL can seamlessly integrate with procedures through its concept of User Defined Functions!
- Below is Scala example


```
val model: LogisticRegressionModel = ...  
  
ctx.udf.register("predict",  
  (x: Float, y: Float) => model.predict(Vector(x, y)))  
  
ctx.sql("SELECT predict(age, weight) FROM users")
```

[2] Armbrust, Michael, et al. "Spark sql: Relational data processing in spark." *Proceedings of the 2015 ACM SIGMOD international conference on management of data*. ACM, 2015



Spark SQL Optimizer Catalyst^[1]

- All the statements are cached as Abstract Syntax Tree (AST)
- Lazy evaluation of AST enables optimization of expressed operations.
- Diagram here depicts optimization pipeline



[2] Armbrust, Michael, et al. "Spark sql: Relational data processing in spark." *Proceedings of the 2015 ACM SIGMOD international conference on management of data*. ACM, 2015



Spark SQL Optimizer Catalyst^[1]

- Optimizer can work based on rules and use some cost base model for choosing a execution plan.
- Catalyst is designed to be Extensible, and workload specific optimizers can be created.

[2] Armbrust, Michael, et al. "Spark sql: Relational data processing in spark." *Proceedings of the 2015 ACM SIGMOD international conference on management of data*. ACM, 2015



References

- [1] Armbrust, Michael, et al. "Spark SQL: Relational data processing in spark." *Proceedings of the 2015 ACM SIGMOD international conference on management of data*. ACM, 2015.
- [2] Spark SQL Getting Started:
<https://spark.apache.org/docs/latest/sql-getting-started.html>
- [3] Spark SQL reference
<https://docs.databricks.com/spark/latest/spark-sql/index.html>
- [4] Documentation pages (version 2.4)
<https://spark.apache.org/docs/2.4.0/>